




Chapter 1: Introduction

- Purpose of Database Systems
- View of Data
- Data Models
- Data Definition Language
- Data Manipulation Language
- Transaction Management
- Storage Management
- Database Administrator
- Database Users
- Overall System Structure
- DBMS Vs. IRS



Database Management System (DBMS)


- Collection of interrelated data
- Set of programs to access the data
- DBMS contains information about a particular enterprise
- DBMS provides an environment that is both *convenient* and *efficient* to use.
- Database Applications:
 - ⌚ Banking: all transactions
 - ⌚ Airlines: reservations, schedules
 - ⌚ Universities: registration, grades
 - ⌚ Sales: customers, products, purchases
 - ⌚ Manufacturing: production, inventory, orders, supply chain
 - ⌚ Human resources: employee records, salaries, tax deductions
- Databases touch all aspects of our lives





Purpose of Database System

- In the early days, database applications were built on top of file systems
- Drawbacks of using file systems to store data:
 - ⌚ Data redundancy and inconsistency
 - ⌚ Multiple file formats, duplication of information in different files
 - ⌚ Difficulty in accessing data
 - ⌚ Need to write a new program to carry out each new task
 - ⌚ Data isolation — multiple files and formats
 - ⌚ Integrity problems
 - ⌚ Integrity constraints (e.g. account balance > 0) become part of program code
 - ⌚ Hard to add new constraints or change existing ones



Purpose of Database Systems (Cont.)

- Drawbacks of using file systems (cont.)
 - ⌚ Atomicity of updates
 - ⌚ Failures may leave database in an inconsistent state with partial updates carried out
 - ⌚ E.g. transfer of funds from one account to another should either complete or not happen at all
 - ⌚ Concurrent access by multiple users
 - ⌚ Concurrent accessed needed for performance
 - ⌚ Uncontrolled concurrent accesses can lead to inconsistencies
 - E.g. two people reading a balance and updating it at the same time
 - ⌚ Security problems
 - Database systems offer solutions to all the above problems





Is the WWW a DBMS?

- Fairly sophisticated search available
 - ❖ crawler indexes pages for fast search
- But, currently
 - ❖ data is mostly unstructured and untyped
 - ❖ can't manipulate the data
 - ❖ few guarantees provided for freshness of data, consistency across data items, fault tolerance, ...
 - ❖ Web sites typically have a DBMS in the background to provide these functions.
- The picture is quickly changing
 - ❖ New standards like XML can help data modeling
 - ❖ Research groups are working on providing some of this functionality across multiple web sites.



Is a File System a DBMS?

- Thought Experiment 1:
 - ❖ You and your project partner are editing the same file.
 - ❖ You both save it at the same time.
 - ❖ Whose changes survive?
A) Yours B) Partner's C) Both D) Neither E) ???
- Thought Experiment 2:
 - ❖ You're updating a file.
 - ❖ The power goes out.
 - ❖ Which of your changes survive?
A) All B) None C) All Since last save D) ???





Levels of Abstraction

- Physical level describes how a record (e.g., customer) is stored.
- Logical level: describes data stored in database, and the relationships among the data.


```
type customer = record
    name : string;
    street : string;
    city : integer;
end;
```

- View level: application programs hide details of data types. Views can also hide information (e.g., salary) for security purposes.



View of Data

An architecture for a database system





Instances and Schemas

- Similar to types and variables in programming languages
- **Schema** – the logical structure of the database
 - ❖ e.g., the database consists of information about a set of customers and accounts and the relationship between them)
 - ❖ Analogous to type information of a variable in a program
 - ❖ **Physical schema:** database design at the physical level
 - ❖ **Logical schema:** database design at the logical level
- **Instance** – the actual content of the database at a particular point in time
 - ❖ Analogous to the value of a variable
- **Physical Data Independence** – the ability to modify the physical schema without changing the logical schema
 - ❖ Applications depend on the logical schema
 - ❖ In general, the interfaces between the various levels and components should be well defined so that changes in some parts do not seriously influence others.



Data Models


- A collection of tools for describing
 - ❖ data
 - ❖ data relationships
 - ❖ data semantics
 - ❖ data constraints
- Entity-Relationship model
- Relational model
- Other models:
 - ❖ object-oriented model
 - ❖ semi-structured data models
- ❖ Older models: network model and hierarchical model





Entity-Relationship Model

Example of schema in the entity-relationship model



Database System Concepts

1.11

©Silberschatz, Korth and Sudarshan



Entity Relationship Model (Cont.)

- E-R model of real world
 - ❖ Entities (objects)
 - ❖ E.g. customers, accounts, bank branch
 - ❖ Relationships between entities
 - ❖ E.g. Account A-101 is held by customer Johnson
 - ❖ Relationship set *depositor* associates customers with accounts
- Widely used for database design
 - ❖ Database design in E-R model usually converted to design in the relational model (coming up next) which is used for storage and processing

Database System Concepts

1.12

©Silberschatz, Korth and Sudarshan



Relational Model

- Example of tabular data in the relational model

<i>Customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>	<i>account-number</i>
192-83-7465	Johnson	Alma	Palo Alto	A-101
019-28-3746	Smith	North	Rye	A-215
192-83-7465	Johnson	Alma	Palo Alto	A-201
321-12-3123	Jones	Main	Harrison	A-217
019-28-3746	Smith	North	Rye	A-201

Attributes

A Sample Relational Database

<i>customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto
019-28-3746	Smith	4 North St.	Rye
677-89-9011	Hayes	3 Main St.	Harrison
182-73-6091	Turner	123 Putnam Ave.	Stamford
321-12-3123	Jones	100 Main St.	Harrison
336-66-9999	Lindsay	175 Park Ave.	Pittsfield
019-28-3746	Smith	72 North St.	Rye

(a) The *customer* table

<i>account-number</i>	<i>balance</i>
A-101	500
A-215	700
A-102	400
A-305	350
A-201	900
A-217	750
A-222	700

(b) The *account* table

<i>customer-id</i>	<i>account-number</i>
192-83-7465	A-101
192-83-7465	A-201
019-28-3746	A-215
677-89-9011	A-102
182-73-6091	A-305
321-12-3123	A-217
336-66-9999	A-222
019-28-3746	A-201


(c) The *depositor* table



Data Definition Language (DDL)

- Specification notation for defining the database schema
 - ❖ E.g.

```
create table account (
    account-number  char(10),
    balance         integer)
```
- DDL compiler generates a set of tables stored in a *data dictionary*
- Data dictionary contains metadata (i.e., data about data)
 - ❖ database schema
 - ❖ Data storage and definition language
 - ❑ language in which the storage structure and access methods used by the database system are specified
 - ❑ Usually an extension of the data definition language



Data Manipulation Language (DML)

- Language for accessing and manipulating the data organized by the appropriate data model
 - ❖ DML also known as query language
- Two classes of languages
 - ❖ Procedural – user specifies what data is required and how to get those data
 - ❖ Nonprocedural – user specifies what data is required without specifying how to get those data
- SQL is the most widely used query language





SQL

- SQL: widely used non-procedural language

⌚ E.g. find the name of the customer with customer-id 192-83-7465

```
select customer.customer-name  
from customer  
where customer.customer-id = '192-83-7465'
```

⌚ E.g. find the balances of all accounts held by the customer with customer-id 192-83-7465

```
select account.balance  
from depositor, account  
where depositor.customer-id = '192-83-7465' and  
depositor.account-number = account.account-number
```

- Application programs generally access databases through


⌚ Language extensions to allow embedded SQL

⌚ Application program interface (e.g. ODBC/JDBC) which allow SQL queries to be sent to a database



Database Users

- Users are differentiated by the way they expect to interact with the system
- Application programmers – interact with system through DML calls
- Sophisticated users – form requests in a database query language
- Specialized users – write specialized database applications that do not fit into the traditional data processing framework
- Naïve users – invoke one of the permanent application programs that have been written previously
 - ⌚ E.g. people accessing database over the web, bank tellers, clerical staff






Database Administrator

- Coordinates all the activities of the database system; the database administrator has a good understanding of the enterprise's information resources and needs.
- Database administrator's duties include:
 - ❖ Schema definition
 - ❖ Storage structure and access method definition
 - ❖ Schema and physical organization modification
 - ❖ Granting user authority to access the database
 - ❖ Specifying integrity constraints
 - ❖ Acting as liaison with users
 - ❖ Monitoring performance and responding to changes in requirements



Transaction Management

- A *transaction* is a collection of operations that performs a single logical function in a database application
- Transaction-management component ensures that the database remains in a consistent (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures.
- Concurrency-control manager controls the interaction among the concurrent transactions, to ensure the consistency of the database.






Storage Management


- Storage manager is a program module that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system.
- The storage manager is responsible to the following tasks:
 - ❖ interaction with the file manager
 - ❖ efficient storing, retrieving and updating of data



Overall System Structure



Application Architectures



▪ **Two-tier architecture:** E.g. client programs using ODBC/JDBC to communicate with a database



▪ **Three-tier architecture:** E.g. web-based applications, and applications built using “middleware”

Advantages of a DBMS

- **Data independence**
- **Efficient data access**
- **Data integrity & security**
- **Data administration**
- **Concurrent access, crash recovery**
- **Reduced application development time**

■ So why not use them always?

- ⌚ Expensive/complicated to set up & maintain
- ⌚ This cost & complexity must be offset by need
- ⌚ General-purpose, not suited for special-purpose tasks (e.g. text search!)





Chapter 2: Entity-Relationship Model


- Entity Sets
- Relationship Sets
- Design Issues
- Mapping Constraints
- Keys
- E-R Diagram
- Extended E-R Features
- Design of an E-R Database Schema
- Reduction of an E-R Schema to Tables



Entity Sets

- A *database* can be modeled as:
 - ❖ a collection of entities,
 - ❖ relationship among entities.
- An *entity* is an object that exists and is distinguishable from other objects.
 - ❖ Example: specific person, company, event, plant
- Entities have *attributes*
 - ❖ Example: people have *names* and *addresses*
- An *entity set* is a set of entities of the same type that share the same properties.
 - ❖ Example: set of all persons, companies, trees, holidays






Database System Concepts

2.3

©Silberschatz, Korth and Sudarshan



Database System Concepts

2.4

©Silberschatz, Korth and Sudarshan

Composite Attributes

Composite Attributes

```

graph TD
    name["name"] --- first_name["first-name"]
    name --- middle_initial["middle-initial"]
    name --- last_name["last-name"]
    address["address"] --- street["street"]
    address --- city["city"]
    address --- state["state"]
    address --- postal_code["postal-code"]
    street --- street_number["street-number"]
    street --- street_name["street-name"]
    street --- apartment_number["apartment-number"]
  
```

Component Attributes

©Silberschatz, Korth and Sudarshan

Relationship Sets

- A **relationship** is an association among several entities

Example:

<u>Hayes</u>	<u>depositor</u>	<u>A-102</u>
customer entity	relationship set	account entity

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets



$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where (e_1, e_2, \dots, e_n) is a relationship

Example:

(Hayes, A-102) \in depositor

©Silberschatz, Korth and Sudarshan






Degree of a Relationship Set



- Refers to number of entity sets that participate in a relationship set.
- Relationship sets that involve two entity sets are *binary* (or degree two). Generally, most relationship sets in a database system are binary.
- Relationship sets may involve more than two entity sets.
 - ❖ E.g. Suppose employees of a bank may have jobs (responsibilities) at multiple branches, with different jobs at different branches. Then there is a ternary relationship set between entity sets *employee, job and branch*
- Relationships between more than two entity sets are rare. Most relationships are binary. (More on this later.)



Mapping Cardinalities

- Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
 - ❖ One to one
 - ❖ One to many
 - ❖ Many to one
 - ❖ Many to many








Mapping Cardinalities affect ER Design

- Can make *access-date* an attribute of account, instead of a relationship attribute, if each account can have only one customer
 - I.e., the relationship from account to customer is many to one, or equivalently, customer to account is one to many




Database System Concepts

2.13


©Silberschatz, Korth and Sudarshan



E-R Diagrams





- Rectangles represent entity sets.
- Diamonds represent relationship sets.
- Lines link attributes to entity sets and entity sets to relationship sets.
- Ellipses represent attributes
 - Double ellipses represent multivalued attributes.
 - Dashed ellipses denote derived attributes.
- Underline indicates primary key attributes (will study later)



Database System Concepts


2.14

©Silberschatz, Korth and Sudarshan



Roles

- Entity sets of a relationship need not be distinct
- The labels “manager” and “worker” are called **roles**; they specify how employee entities interact via the works-for relationship set.
- Roles are indicated in E-R diagrams by labeling the lines that connect diamonds to rectangles.
- Role labels are optional, and are used to clarify semantics of the relationship




Database System Concepts

2.17

©Silberschatz, Korth and Sudarshan

Cardinality Constraints

- We express cardinality constraints by drawing either a directed line (\rightarrow), signifying “one,” or an undirected line ($-$), signifying “many,” between the relationship set and the entity set.
- E.g.: One-to-one relationship:
 - ⌚ A customer is associated with at most one loan via the relationship *borrower*
 - ⌚ A loan is associated with at most one customer via *borrower*



Database System Concepts


2.18

©Silberschatz, Korth and Sudarshan



One-To-Many Relationship

- In the one-to-many relationship a loan is associated with at most one customer via *borrower*, a customer is associated with several (including 0) loans via *borrower*



Database System Concepts


2.19

©Silberschatz, Korth and Sudarshan



Many-To-One Relationships

- In a many-to-one relationship a loan is associated with several (including 0) customers via *borrower*, a customer is associated with at most one loan via *borrower*



Database System Concepts

2.20

©Silberschatz, Korth and Sudarshan



Many-To-Many Relationship

The ER diagram illustrates a many-to-many relationship between the *customer* and *loan* entity sets. The *borrower* relationship set connects them. The *customer* entity set contains attributes *customer-name*, *customer-street*, *customer-id*, and *customer-city*. The *loan* entity set contains attributes *loan-number* and *amount*. The *borrower* relationship set is represented by a diamond shape connecting the two entity sets.

- A customer is associated with several (possibly 0) loans via borrower
- A loan is associated with several (possibly 0) customers via borrower

Database System Concepts 2.21 ©Silberschatz, Korth and Sudarshan

Participation of an Entity Set in a Relationship Set

The ER diagram illustrates the participation of the *customer* entity set in the *borrower* relationship set. Total participation is indicated by a double line on the relationship side. The *customer* entity set contains attributes *customer-name*, *customer-street*, *customer-id*, and *customer-city*. The *loan* entity set contains attributes *loan-number* and *amount*. The *borrower* relationship set is represented by a diamond shape connecting the two entity sets.


- **Total participation** (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
 - E.g. participation of *loan* in *borrower* is total
 - every loan must have a customer associated to it via borrower
- **Partial participation**: some entities may not participate in any relationship in the relationship set
 - E.g. participation of *customer* in *borrower* is partial

Database System Concepts 2.22 ©Silberschatz, Korth and Sudarshan



Alternative Notation for Cardinality Limits

- Cardinality limits can also express participation constraints



Database System Concepts

2.23

©Silberschatz, Korth and Sudarshan



Keys

- A **super key** of an entity set is a set of one or more attributes whose values uniquely determine each entity.
- A **candidate key** of an entity set is a minimal super key
 - Customer-id* is candidate key of *customer*
 - account-number* is candidate key of *account*
- Although several candidate keys may exist, one of the candidate keys is selected to be the **primary key**.

Database System Concepts

2.24

©Silberschatz, Korth and Sudarshan







Keys for Relationship Sets

- The combination of primary keys of the participating entity sets forms a super key of a relationship set.
 - ❖ $(customer-id, account-number)$ is the super key of *depositor*
 - ❖ NOTE: *this means a pair of entity sets can have at most one relationship in a particular relationship set.*
 - ▀ E.g. if we wish to track all access-dates to each account by each customer, we cannot assume a relationship for each access. We can use a multivalued attribute though
- Must consider the mapping cardinality of the relationship set when deciding the what are the candidate keys
- Need to consider semantics of relationship set in selecting the *primary key* in case of more than one candidate key




E-R Diagram with a Ternary Relationship





Cardinality Constraints on Ternary Relationship

- We allow at most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint
- E.g. an arrow from *works-on* to *job* indicates each employee works on at most one job at any branch.
- If there is more than one arrow, there are two ways of defining the meaning.
 - ⌚ E.g. a ternary relationship R between A , B and C with arrows to B and C could mean
 1. each A entity is associated with a unique entity from B and C or
 2. each pair of entities from (A, B) is associated with a unique C entity, and each pair (A, C) is associated with a unique B
 - ⌚ Each alternative has been used in different formalisms
 - ⌚ To avoid confusion we outlaw more than one arrow




Binary Vs. Non-Binary Relationships

- Some relationships that appear to be non-binary may be better represented using binary relationships
 - ⌚ E.g. A ternary relationship *parents*, relating a child to his/her father and mother, is best replaced by two binary relationships, *father* and *mother*
 - ⌚ Using two binary relationships allows partial information (e.g. only mother being known)
 - ⌚ But there are some relationships that are naturally non-binary
 - ⌚ E.g. *works-on*



Converting Non-Binary Relationships to Binary Form

- In general, any non-binary relationship can be represented using binary relationships by creating an artificial entity set E , and three relationship sets:
 - ❖ Replace R between entity sets A , B and C by an entity set E , and three relationship sets:
 1. R_A , relating E and A
 2. R_B , relating E and B
 3. R_C , relating E and C
 - ❖ Create a special identifying attribute for E
 - ❖ Add any attributes of R to E
 - ❖ For each relationship (a_i, b_i, c_i) in R , create
 1. a new entity e_i in the entity set E
 2. add (e_i, a_i) to R_A
 3. add (e_i, b_i) to R_B
 4. add (e_i, c_i) to R_C



Database System Concepts

2.29

©Silberschatz, Korth and Sudarshan

Converting Non-Binary Relationships (Cont.)

- Also need to translate constraints
 - ❖ Translating all constraints may not be possible
 - ❖ There may be instances in the translated schema that cannot correspond to any instance of R
 - *Exercise: add constraints to the relationships R_A , R_B and R_C to ensure that a newly created entity corresponds to exactly one entity in each of entity sets A , B and C*
 - ❖ We can avoid creating an identifying attribute by making E a weak entity set (described shortly) identified by the three relationship sets

Database System Concepts

2.30

©Silberschatz, Korth and Sudarshan


Weak Entity Sets

- An entity set that does not have a primary key is referred to as a *weak entity set*.
- The existence of a weak entity set depends on the existence of a *identifying entity set*
 - ❖ it must relate to the identifying entity set via a total, one-to-many relationship set from the identifying to the weak entity set
 - ❖ *Identifying relationship* depicted using a double diamond
- The *discriminator* (or *partial key*) of a weak entity set is the set of attributes that distinguishes among all the entities of a weak entity set.
- The primary key of a weak entity set is formed by the primary key of the strong entity set on which the weak entity set is existence dependent, plus the weak entity set's discriminator.



Weak Entity Sets (Cont.)

- We depict a weak entity set by double rectangles.
- We underline the discriminator of a weak entity set with a dashed line.
- *payment-number* – discriminator of the *payment* entity set
- Primary key for *payment* – (*loan-number*, *payment-number*)





Weak Entity Sets (Cont.)


- Note: the primary key of the strong entity set is not explicitly stored with the weak entity set, since it is implicit in the identifying relationship.
- If *loan-number* were explicitly stored, *payment* could be made a strong entity, but then the relationship between *payment* and *loan* would be duplicated by an implicit relationship defined by the attribute *loan-number* common to *payment* and *loan*



Example: Logins (Email Addresses)



Login name = user name + host name, e.g.,
ark@soe.ucsc.edu.



- A “login” entity corresponds to a user name on a particular host, but the passwd table doesn’t record the host, just the user name, e.g., ark.
- Key for a login = the user name at the host (which is unique for that host only) + the IP address of the host (which is unique globally).



- Design issue: Under what circumstances could we simply make *login-name* and *host-name* be attributes of *logins*, and dispense with the weak E.S.?










Entity Sets Vs. Attributes

You may be unsure which concepts are worthy of being entity sets, and which are handled more simply as attributes.

- Especially tricky for the class design project, since there is a temptation to create needless entity sets to make project "larger."




Make an entity set only if it either:
 Is more than a name of something; i.e., it has nonkey attributes or relationships with a number of different entity sets,
 or
 Is the "many" in a many-one relationship.





Example

The following design illustrates both points:



- Manfs* deserves to be an E.S. because we record *addr*, a nonkey attribute.
- Beers* deserves to be an E.S. because it is at the "many" end.
 - If not, we would have to make "set of beers" an attribute of *Manfs* – something we avoid doing, although some may tell you it is OK in E/R model.







How about doing an ER design
interactively on the board?
Suggest an application to be modeled.

Specialization

- Top-down design process; we designate subgroupings within an entity set that are distinctive from other entities in the set.
- These subgroupings become lower-level entity sets that have attributes or participate in relationships that do not apply to the higher-level entity set.
- Depicted by a *triangle* component labeled ISA (E.g. *customer* “is a” *person*).
- **Attribute inheritance** – a lower-level entity set inherits all the attributes and relationship participation of the higher-level entity set to which it is linked.









Specialization and Generalization (Contd.)

- Can have multiple specializations of an entity set based on different features.
- E.g. *permanent-employee* vs. *temporary-employee*, in addition to *officer* vs. *secretary* vs. *teller*
- Each particular employee would be
 - ❖ a member of one of *permanent-employee* or *temporary-employee*,
 - ❖ and also a member of one of *officer*, *secretary*, or *teller*
- The ISA relationship also referred to as **superclass - subclass** relationship



Design Constraints on a Specialization/Generalization

- Constraint on which entities can be members of a given lower-level entity set.
 - ❖ condition-defined
 - ▀ E.g. all customers over 65 years are members of *senior-citizen* entity set; *senior-citizen* ISA *person*.
 - ❖ user-defined
- Constraint on whether or not entities may belong to more than one lower-level entity set within a single generalization.
 - ❖ **Disjoint**
 - ▀ an entity can belong to only one lower-level entity set
 - ▀ Noted in E-R diagram by writing *disjoint* next to the ISA triangle
 - ❖ **Overlapping**
 - ▀ an entity can belong to more than one lower-level entity set






Design Constraints on a Specialization/Generalization (Contd.)

- **Completeness constraint** -- specifies whether or not an entity in the higher-level entity set must belong to at least one of the lower-level entity sets within a generalization.
 - ❖ **total** : an entity must belong to one of the lower-level entity sets
 - ❖ **partial**: an entity need not belong to one of the lower-level entity sets



Aggregation

- Consider the ternary relationship *works-on*, which we saw earlier
- Suppose we want to record managers for tasks performed by an employee at a branch




Aggregation (Cont.)

- Relationship sets *works-on* and *manages* represent overlapping information
 - ❖ Every *manages* relationship corresponds to a *works-on* relationship
 - ❖ However, some *works-on* relationships may not correspond to any *manages* relationships
 - ▀ So we can't discard the *works-on* relationship
- Eliminate this redundancy via *aggregation*
 - ❖ Treat relationship as an abstract entity
 - ❖ Allows relationships between relationships
 - ❖ Abstraction of relationship into new entity
- Without introducing redundancy, the following diagram represents:
 - ❖ An employee works on a particular job at a particular branch
 - ❖ An employee, branch, job combination may have an associated manager




E-R Diagram With Aggregation








E-R Design Decisions

- The use of an attribute or entity set to represent an object.
- Whether a real-world concept is best expressed by an entity set or a relationship set.
- The use of a ternary relationship versus a pair of binary relationships.
- The use of a strong or weak entity set.
- The use of specialization/generalization – contributes to modularity in the design.
- The use of aggregation – can treat the aggregate entity set as a single unit without concern for the details of its internal structure.



Beers-Bars-Drinkers Example









Database System Concepts



2.55

©Silberschatz, Korth and Sudarshan

How about doing another ER design
interactively on the board?










UML Class Diagram Notation (Cont.)

3. Cardinality constraints



4. Generalization and Specialization




*Note reversal of position in cardinality constraint depiction

*Generalization can use merged or separate arrows independent of disjoint/overlapping



UML Class Diagrams (Contd.)


- Cardinality constraints are specified in the form $l..h$, where l denotes the minimum and h the maximum number of relationships an entity can participate in.
- Beware: the positioning of the constraints is exactly the reverse of the positioning of constraints in E-R diagrams.
- The constraint $0..*$ on the $E2$ side and $0..1$ on the $E1$ side means that each $E2$ entity can participate in at most one relationship, whereas each $E1$ entity can participate in many relationships; in other words, the relationship is many to one from $E2$ to $E1$.
- Single values, such as 1 or * may be written on edges; The single value 1 on an edge is treated as equivalent to $1..1$, while * is equivalent to $0..*$.





Reduction of an E-R Schema to Tables

- Primary keys allow entity sets and relationship sets to be expressed uniformly as *tables* which represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of tables.
- For each entity set and relationship set there is a unique table which is assigned the name of the corresponding entity set or relationship set.
- Each table has a number of columns (generally corresponding to attributes), which have unique names.
- Converting an E-R diagram to a table format is the basis for deriving a relational database design from an E-R diagram.



Representing Entity Sets as Tables

- A strong entity set reduces to a table with the same attributes.

<i>customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
019-28-3746	Smith	North	Rye
182-73-6091	Turner	Putnam	Stamford
192-83-7465	Johnson	Alma	Palo Alto
244-66-8800	Curry	North	Rye
321-12-3123	Jones	Main	Harrison
335-57-7991	Adams	Spring	Pittsfield
336-66-9999	Lindsay	Park	Pittsfield
677-89-9011	Hayes	Main	Harrison
963-96-3963	Williams	Nassau	Princeton





Composite and Multivalued Attributes

- Composite attributes are flattened out by creating a separate attribute for each component attribute
 - ❖ E.g. given entity set *customer* with composite attribute *name* with component attributes *first-name* and *last-name* the table corresponding to the entity set has two attributes *name.first-name* and *name.last-name*
- A multivalued attribute M of an entity E is represented by a separate table EM
 - ❖ Table EM has attributes corresponding to the primary key of E and an attribute corresponding to multivalued attribute M
 - ❖ E.g. Multivalued attribute *dependent-names* of *employee* is represented by a table *employee-dependent-names(employee-id, dname)*
 - ❖ Each value of the multivalued attribute maps to a separate row of the table EM
 - ▀ E.g., an employee entity with primary key John and dependents Johnson and Johndotir maps to two rows: (John, Johnson) and (John, Johndotir)



Representing Weak Entity Sets

- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set

<i>loan-number</i>	<i>payment-number</i>	<i>payment-date</i>	<i>payment-amount</i>
L-11	53	7 June 2001	125
L-14	69	28 May 2001	500
L-15	22	23 May 2001	300
L-16	58	18 June 2001	135
L-17	5	10 May 2001	50
L-17	6	7 June 2001	50
L-17	7	17 June 2001	100
L-23	11	17 May 2001	75
L-93	103	3 June 2001	900
L-93	104	13 June 2001	200



Representing Relationship Sets as Tables

- A many-to-many relationship set is represented as a table with columns for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- E.g.: table for relationship set *borrower*

<i>customer-id</i>	<i>loan-number</i>
019-28-3746	L-11
019-28-3746	L-23
244-66-8800	L-93
321-12-3123	L-17
335-57-7991	L-16
555-55-5555	L-14
677-89-9011	L-15
963-96-3963	L-17





Database System Concepts

2.69

©Silberschatz, Korth and Sudarshan

Redundancy of Tables

- Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the many side, containing the primary key of the one side
- E.g.: Instead of creating a table for relationship *account-branch*, add an attribute *branch* to the entity set *account*



Database System Concepts

2.70

©Silberschatz, Korth and Sudarshan



Redundancy of Tables (Cont.)

- For one-to-one relationship sets, either side can be chosen to act as the “many” side
 - ❖ That is, extra attribute can be added to either of the tables corresponding to the two entity sets
- If participation is *partial* on the many side, replacing a table by an extra attribute in the relation corresponding to the “many” side could result in null values
- The table corresponding to a relationship set linking a weak entity set to its identifying strong entity set is redundant.
 - ❖ E.g. The *payment* table already contains the information that would appear in the *loan-payment* table (i.e., the columns *loan-number* and *payment-number*).



Representing Specialization as Tables

■ Method 1:

- ❖ Form a table for the higher level entity
- ❖ Form a table for each lower level entity set, include primary key of higher level entity set and local attributes

table	table attributes
<i>person</i>	<i>name, street, city</i>
<i>customer</i>	<i>name, credit-rating</i>
<i>employee</i>	<i>name, salary</i>

- ❖ Drawback: getting information about, e.g., *employee* requires accessing two tables





Representing Specialization as Tables (Cont.)

■ Method 2:

- ❖ Form a table for each entity set with all local and inherited attributes

table	table attributes
<i>person</i>	<i>name, street, city</i>
<i>customer</i>	<i>name, street, city, credit-rating</i>
<i>employee</i>	<i>name, street, city, salary</i>

- ❖ If specialization is total, table for generalized entity (*person*) not required to store information
 - ❖ Can be defined as a “view” relation containing union of specialization tables
 - ❖ But explicit table may still be needed for foreign key constraints
- ❖ Drawback: street and city may be stored redundantly for persons who are both customers and employees



Relations Corresponding to Aggregation


- To represent aggregation, create a table containing
 - primary key of the aggregated relationship,
 - the primary key of the associated entity set
 - Any descriptive attributes





Relations Corresponding to Aggregation (Cont.)

- E.g. to represent aggregation *manages* between relationship *works-on* and entity set *manager*, create a table *manages(employee-id, branch-name, title, manager-name)*
- Table *works-on* is redundant **provided** we are willing to store null values for attribute *manager-name* in table *manages*





Database System Concepts


2.75

©Silberschatz, Korth and Sudarshan



End of Chapter 2






Database System Concepts

2.79

©Silberschatz, Korth and Sudarshan



Database System Concepts


2.80

©Silberschatz, Korth and Sudarshan



Existence Dependencies

- If the existence of entity x depends on the existence of entity y , then x is said to be *existence dependent* on y .
 - ⌚ y is a *dominant entity* (in example below, *loan*)
 - ⌚ x is a *subordinate entity* (in example below, *payment*)




If a *loan* entity is deleted, then all its associated *payment* entities must be deleted also.





Chapter 3: Relational Model

- Structure of Relational Databases
- Relational Algebra
- Tuple Relational Calculus
- Domain Relational Calculus
- Extended Relational-Algebra-Operations
- Modification of the Database
- Views



Example of a Relation

account-number	branch-name	balance
A-101	Downtown	500
A-102	Perryridge	400
A-201	Brighton	900
A-215	Mianus	700
A-217	Brighton	750
A-222	Redwood	700
A-305	Round Hill	350





Basic Structure

- Formally, given sets D_1, D_2, \dots, D_n a **relation** r is a subset of $D_1 \times D_2 \times \dots \times D_n$
Thus a relation is a set of n-tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$
- Example: if

customer-name = {Jones, Smith, Curry, Lindsay}

customer-street = {Main, North, Park}

customer-city = {Harrison, Rye, Pittsfield}


Then $r = \{ (Jones, Main, Harrison), (Smith, North, Rye), (Curry, North, Rye), (Lindsay, Park, Pittsfield) \}$

is a relation over *customer-name* \times *customer-street* \times *customer-city*



Attribute Types


- Each attribute of a relation has a name
- The set of allowed values for each attribute is called the **domain** of the attribute
- Attribute values are (normally) required to be **atomic**, that is, indivisible
 - ⌚ E.g. multivalued attribute values are not atomic
 - ⌚ E.g. composite attribute values are not atomic
- The special value *null* is a member of every domain
- The null value causes complications in the definition of many operations
 - ⌚ we shall ignore the effect of null values in our main presentation and consider their effect later





Relation Schema

- A_1, A_2, \dots, A_n are *attributes*
- $R = (A_1, A_2, \dots, A_n)$ is a *relation schema*
 - E.g. *Customer-schema* =
 $(\text{customer-name}, \text{customer-street}, \text{customer-city})$
- $r(R)$ is a *relation* on the *relation schema* R
 - E.g. *customer* (*Customer-schema*)



Relation Instance


- The current values (*relation instance*) of a relation are specified by a table
- An element t of r is a *tuple*, represented by a *row* in a table

<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
Jones	Main	Harrison
Smith	North	Rye
Curry	North	Rye
Lindsay	Park	Pittsfield

customer

Diagram illustrating the components of a relation instance:

- The table has three columns labeled *customer-name*, *customer-street*, and *customer-city*.
- The rows represent tuples (or rows), with arrows pointing from the column labels to the first row.
- Labels "attributes (or columns)" point to the column headers, and "tuples (or rows)" point to the data rows.





Relations are Unordered


- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
- E.g. *account* relation with unordered tuples

account-number	branch-name	balance
A-101	Downtown	500
A-215	Mianus	700
A-102	Perryridge	400
A-305	Round Hill	350
A-201	Brighton	900
A-222	Redwood	700
A-217	Brighton	750



Why Relations?

- Very simple model.
- Often a good match for the way we think about our data.
- Abstract model that underlies SQL, the most important language in DBMS's today.
 - ⌚ But SQL uses "bags" while the abstract relational model is set-oriented.
- All ingenious ideas are simple !





Database

- A database consists of multiple relations
- Information about an enterprise is broken up into parts, with each relation storing one part of the information

E.g.: *account*: stores information about accounts
depositor: stores information about which customer owns which account
customer: stores information about customers

- Storing all information as a single relation such as *bank(account-number, balance, customer-name, ..)* results in
 - ⌚ repetition of information (e.g. two customers own an account)
 - ⌚ the need for null values (e.g. represent a customer without an account)
- Normalization theory (Chapter 7) deals with how to design relational schemas



The *customer* Relation

<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
Adams	Spring	Pittsfield
Brooks	Senator	Brooklyn
Curry	North	Rye
Glenn	Sand Hill	Woodside
Green	Walnut	Stamford
Hayes	Main	Harrison
Johnson	Alma	Palo Alto
Jones	Main	Harrison
Lindsay	Park	Pittsfield
Smith	North	Rye
Turner	Putnam	Stamford
Williams	Nassau	Princeton



The depositor Relation


<i>customer-name</i>	<i>account-number</i>
Hayes	A-102
Johnson	A-101
Johnson	A-201
Jones	A-217
Lindsay	A-222
Smith	A-215
Turner	A-305

Database System Concepts

3.11

©Silberschatz, Korth and Sudarshan

E-R Diagram for the Banking Enterprise



Database System Concepts

3.12

©Silberschatz, Korth and Sudarshan



Keys

- Let $K \subseteq R$
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
 - ❖ by “possible r ” we mean a relation r that could exist in the enterprise we are modeling.
 - ❖ Example: $\{customer-name, customer-street\}$ and $\{customer-name\}$ are both superkeys of *Customer*, if no two customers can possibly have the same name.
- K is a **candidate key** if K is minimal
 - Example: $\{customer-name\}$ is a candidate key for *Customer*, since it is a superkey (assuming no two customers can possibly have the same name), and no subset of it is a superkey.




Example 1

Drinkers(name, addr, beersLiked, manf, favoriteBeer)

- $\{\text{name}, \text{beersLiked}\}$ FD's all attributes, as seen.
 - ❖ Shows $\{\text{name}, \text{beersLiked}\}$ is a superkey.
- $\text{name} \rightarrow \text{beersLiked}$ is false, so name is not a superkey.
- $\text{beersLiked} \rightarrow \text{name}$ also false, so beersLiked is not a superkey.
- Thus, $\{\text{name}, \text{beersLiked}\}$ is a key.
 - ❖ Neither name nor beersLiked is on the right of any observed FD, so they must be part of *any* superkey.
- Important point: “key” in a relation refers to tuples, not the entities they represent. If an entity is represented by several tuples, then entity-key will not be the same as relation-key.



Example 2




Note: There are alternate keys

- Keys are {Lastname, Firstname} and {StudentID}




Determining Keys from E-R Sets

- **Strong entity set.** The primary key of the entity set becomes the primary key of the relation.
- **Weak entity set.** The primary key of the relation consists of the union of the primary key of the strong entity set and the discriminator of the weak entity set.
- **Relationship set.** The union of the primary keys of the related entity sets becomes a super key of the relation.
 - ❖ For binary many-to-one relationship sets, the primary key of the "many" entity set becomes the relation's primary key.
 - ❖ For one-to-one relationship sets, the relation's primary key can be that of either entity set.
 - ❖ For many-to-many relationship sets, the union of the primary keys becomes the relation's primary key






Schema Diagram for the Banking Enterprise



Database System Concepts

3.17

©Silberschatz, Korth and Sudarshan



Query Languages

- Language in which user requests information from the database.
- Categories of languages
 - ❖ procedural
 - ❖ non-procedural
- “Pure” languages:
 - ❖ Relational Algebra
 - ❖ Tuple Relational Calculus
 - ❖ Domain Relational Calculus
- Pure languages form underlying basis of query languages that people use.



Database System Concepts


3.18

©Silberschatz, Korth and Sudarshan



Relational Algebra

- Procedural language
- Six basic operators
 - ❖ select
 - ❖ project
 - ❖ union
 - ❖ set difference
 - ❖ Cartesian product
 - ❖ rename
- The operators take one or more relations as inputs and give a new relation as a result.




Select Operation – Example

- Relation r

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

- $\sigma_{A=B \wedge D > 5}(r)$

A	B	C	D
α	α	1	7
β	β	23	10



Select Operation

- Notation: $\sigma_p(r)$
- p is called the **selection predicate**
- Defined as:

$$\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$$


Where p is a formula in propositional calculus consisting of **terms** connected by : \wedge (**and**), \vee (**or**), \neg (**not**)
Each **term** is one of:

<attribute> op <attribute> or <constant>

where op is one of: $=, \neq, >, \geq, <, \leq$

- Example of selection:

$$\sigma_{\text{branch-name}=\text{"Perryridge}}(\text{account})$$



Project Operation – Example

- Relation r


	A	B	C
α	10	1	
α	20	1	
β	30	1	
β	40	2	

- $\Pi_{A,C}(r)$

A	C
α	1
α	1
β	1
β	2

=

A	C
α	1
β	1
β	2






Project Operation

- Notation:

$$\Pi_{A_1, A_2, \dots, A_k}(r)$$

where A_1, A_2 are attribute names and r is a relation name.

- The result is defined as the relation of k columns obtained by erasing the columns that are not listed
- Duplicate rows removed from result, since relations are sets
- E.g. To eliminate the *branch-name* attribute of *account*
 $\Pi_{\text{account-number}, \text{balance}}(\text{account})$



Union Operation – Example

- Relations r, s :

A	B
α	1
α	2
β	1


r

A	B
α	2
β	3

s

$r \cup s$:

A	B
α	1
α	2
β	1
β	3





Union Operation

- Notation: $r \cup s$
- Defined as:
$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$
- For $r \cup s$ to be valid.
 1. r, s must have the *same arity* (same number of attributes)
 2. The attribute domains must be *compatible* (e.g., 2nd column of r deals with the same type of values as does the 2nd column of s)
- E.g. to find all customers with either an account or a loan
$$\Pi_{customer-name} (depositor) \cup \Pi_{customer-name} (borrower)$$



Set Difference Operation – Example

- Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

$r - s$:


A	B
α	1
β	1





Set Difference Operation

- Notation $r - s$
- Defined as:
$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$
- Set differences must be taken between *compatible* relations.
 - ⌚ r and s must have the *same arity*
 - ⌚ attribute domains of r and s must be compatible



Cartesian-Product Operation-Example

Relations r, s :

A	B
α	1
β	2


r

C	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

s

$r \times s$:

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b





Cartesian-Product Operation

- Notation $r \times s$
- Defined as:

$$r \times s = \{t q \mid t \in r \text{ and } q \in s\}$$

- Assume that attributes of $r(R)$ and $s(S)$ are disjoint.
(That is, $R \cap S = \emptyset$).
- If attributes of $r(R)$ and $s(S)$ are not disjoint, then renaming must be used.




Composition of Operations

- Can build expressions using multiple operations
- Example: $\sigma_{A=C}(r \times s)$
- $r \times s$

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

- $\sigma_{A=C}(r \times s)$

A	B	C	D	E
α	1	α	10	a
β	2	β	20	a
β	2	β	20	b





Rename Operation

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- Allows us to refer to a relation by more than one name.

Example:


$$\rho_X(E)$$

returns the expression E under the name X

If a relational-algebra expression E has arity n , then

$$\rho_{X(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .



Banking Example

branch (branch-name, branch-city, assets)


customer (customer-name, customer-street, customer-only)

account (account-number, branch-name, balance)

loan (loan-number, branch-name, amount)

depositor (customer-name, account-number)

borrower (customer-name, loan-number)





Example Queries

- Find all loans of over \$1200

$$\sigma_{amount > 1200} (loan)$$

- Find the loan number for each loan of an amount greater than \$1200

$$\Pi_{loan-number} (\sigma_{amount > 1200} (loan))$$


Example Queries

- Find the names of all customers who have a loan, an account, or both, from the bank

$$\Pi_{customer-name} (borrower) \cup \Pi_{customer-name} (depositor)$$

- Find the names of all customers who have a loan and an account at bank.

$$\Pi_{customer-name} (borrower) \cap \Pi_{customer-name} (depositor)$$





Example Queries

- Find the names of all customers who have a loan at the Perryridge branch.

$$\Pi_{customer-name} (\sigma_{branch-name = "Perryridge"} (\sigma_{borrower.loan-number = loan.loan-number} (borrower \times loan)))$$

- Find the names of all customers who have a loan at the Perryridge branch but do not have an account at any branch of the bank.

$$\Pi_{customer-name} (\sigma_{branch-name = "Perryridge"} (\sigma_{borrower.loan-number = loan.loan-number} (borrower \times loan)) - \Pi_{customer-name} (depositor))$$



Example Queries

- Find the names of all customers who have a loan at the Perryridge branch.

– Query 1

$$\Pi_{customer-name} (\sigma_{branch-name = "Perryridge"} (\sigma_{borrower.loan-number = loan.loan-number} (borrower \times loan)))$$

– Query 2

$$\Pi_{customer-name} (\sigma_{loan.loan-number = borrower.loan-number} (\sigma_{branch-name = "Perryridge"} (loan) \times borrower))$$





Example Queries

Find the largest account balance

- Rename *account* relation as *d*
- The query is:

$$\begin{aligned}\Pi_{balance}(account) - \Pi_{account.balance} \\ (\sigma_{account.balance < d.balance} (account \times \rho_d (account)))\end{aligned}$$



Formal Definition

- A basic expression in the relational algebra consists of either one of the following:
 - ❖ A relation in the database
 - ❖ A constant relation
- Let E_1 and E_2 be relational-algebra expressions; the following are all relational-algebra expressions:
 - ❖ $E_1 \cup E_2$
 - ❖ $E_1 - E_2$
 - ❖ $E_1 \times E_2$
 - ❖ $\sigma_p (E_1)$, P is a predicate on attributes in E_1
 - ❖ $\Pi_S(E_1)$, S is a list consisting of some of the attributes in E_1
 - ❖ $\rho_x(E_1)$, x is the new name for the result of E_1






Additional Operations


We define additional operations that do not add any power to the relational algebra, but that simplify common queries.

- Set intersection
- Natural join
- Division
- Assignment



Set-Intersection Operation

- Notation: $r \cap s$
- Defined as:
- $r \cap s = \{ t \mid t \in r \text{ and } t \in s \}$
- Assume:
 - ◊ r, s have the same arity
 - ◊ attributes of r and s are compatible
- Note: $r \cap s = r - (r - s)$





Set-Intersection Operation - Example

- Relation r, s:

A	B
α	1
α	2
β	1


r

A	B
α	2
β	3

s

- $r \cap s$

A	B
α	2



Natural-Join Operation

- Notation: $r \bowtie s$
- Let r and s be relations on schemas R and S respectively. Then, $r \bowtie s$ is a relation on schema $R \cup S$ obtained as follows:

- ❖ Consider each pair of tuples t_r from r and t_s from s .
- ❖ If t_r and t_s have the same value on each of the attributes in $R \cap S$, add a tuple t to the result, where
 - ▀ t has the same value as t_r on r
 - ▀ t has the same value as t_s on s

- Example:

$$R = (A, B, C, D)$$

$$S = (E, B, D)$$

❖ Result schema = (A, B, C, D, E)

❖ $r \bowtie s$ is defined as:

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$$





Natural Join Operation – Example

- Relations r, s:

A	B	C	D
α	1	α	a
β	2	γ	a
γ	4	β	b
α	1	γ	a
δ	2	β	b

r

B	D	E
1	a	α
3	a	β
1	a	γ
2	b	δ
3	b	ϵ

s

$r \bowtie s$

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ



Division Operation

$r \div s$


- Suited to queries that include the phrase “for all”.
- Let r and s be relations on schemas R and S respectively where

- $R = (A_1, \dots, A_m, B_1, \dots, B_n)$
- $S = (B_1, \dots, B_n)$

The result of $r \div s$ is a relation on schema

$$R - S = (A_1, \dots, A_m)$$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r) \}$$





Division Operation – Example

Relations r, s :

	A	B
α	1	
α	2	
α	3	
β	1	
γ	1	
δ	1	
δ	3	
δ	4	
ϵ	6	
ϵ	1	
β	2	

B
1
2

$r \div s$:

A
α
β

r



Another Division Example

Relations r, s :

	A	B	C	D	E
α	a	α	a		1
α	a	γ	a		1
α	a	γ	b		1
β	a	γ	a		1
β	a	γ	b		3
γ	a	γ	a		1
γ	a	γ	b		1
γ	a	β	b		1

D	E
a	1
b	1

s

r

$r \div s$:

A	B	C
α	a	γ
γ	a	γ





Division Operation (Cont.)

- Property

- ❖ Let $q = r \div s$
 - ❖ Then q is the largest relation satisfying $q \times s \subseteq r$


- Definition in terms of the basic algebra operation

Let $r(R)$ and $s(S)$ be relations, and let $S \subseteq R$

$$r \div s = \Pi_{R-S}(r) - \Pi_{R-S}((\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r))$$

To see why

- ❖ $\Pi_{R-S,S}(r)$ simply reorders attributes of r
 - ❖ $\Pi_{R-S}(\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r)$ gives those tuples t in $\Pi_{R-S}(r)$ such that for some tuple $u \in s$, $tu \notin r$.



Assignment Operation


- The assignment operation (\leftarrow) provides a convenient way to express complex queries.

- ❖ Write query as a sequential program consisting of
 - ▀ a series of assignments
 - ▀ followed by an expression whose value is displayed as a result of the query.
 - ❖ Assignment must always be made to a temporary relation variable.

- Example: Write $r \div s$ as

```
temp1  $\leftarrow \Pi_{R-S}(r)
temp2  $\leftarrow \Pi_{R-S}((temp1 \times s) - \Pi_{R-S,S}(r))
result = temp1 - temp2$$ 
```

- ❖ The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow .
 - ❖ May use variable in subsequent expressions.





Example Queries

- Find all customers who have an account from at least the “Downtown” and the Uptown” branches.


Query 1

$$\begin{aligned}\Pi_{CN}(\sigma_{BN=\text{"Downtown"}}(depositor \bowtie account)) \cap \\ \Pi_{CN}(\sigma_{BN=\text{"Uptown"}}(depositor \bowtie account))\end{aligned}$$

where CN denotes customer-name and BN denotes branch-name.

Query 2


$$\begin{aligned}\Pi_{customer-name, branch-name}(depositor \bowtie account) \\ \div \rho_{temp(branch-name)}(\{("Downtown"), ("Uptown")\})\end{aligned}$$



Example Queries

- Find all customers who have an account at all branches located in Brooklyn city.

$$\begin{aligned}\Pi_{customer-name, branch-name}(depositor \bowtie account) \\ \div \Pi_{branch-name}(\sigma_{branch-city = \text{"Brooklyn"}}(branch))\end{aligned}$$





Extended Relational-Algebra-Operations

- Generalized Projection
- Outer Join
- Aggregate Functions



Generalized Projection

- Extends the projection operation by allowing arithmetic functions to be used in the projection list.

$$\Pi_{F_1, F_2, \dots, F_n}(E)$$

- E is any relational-algebra expression
- Each of F_1, F_2, \dots, F_n are arithmetic expressions involving constants and attributes in the schema of E .
- Given relation $credit\text{-}info(customer\text{-}name, limit, credit\text{-}balance)$, find how much more each person can spend:

$$\Pi_{customer\text{-}name, limit - credit\text{-}balance}(credit\text{-}info)$$





Aggregate Functions and Operations

- **Aggregation function** takes a collection of values and returns a single value as a result.

avg: average value
min: minimum value
max: maximum value
sum: sum of values
count: number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \mathcal{G}_{F_1(A_1), F_2(A_2), \dots, F_n(A_n)}(E)$$

- ⌚ E is any relational-algebra expression
- ⌚ G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
- ⌚ Each F_i is an aggregate function
- ⌚ Each A_i is an attribute name



Aggregate Operation – Example

- Relation r :

A	B	C
α	α	7
α	β	7
β	β	3
β	β	10

$$\mathcal{G}_{\text{sum}(C)}(r)$$

sum-C
27






Aggregate Operation – Example

- Relation *account* grouped by *branch-name*:

branch-name	account-number	balance
Perryridge	A-102	400
Perryridge	A-201	900
Brighton	A-217	750
Brighton	A-215	750
Redwood	A-222	700

branch-name $\text{g}_{\text{sum(balance)}}(\text{account})$


branch-name	balance
Perryridge	1300
Brighton	1500
Redwood	700



Aggregate Functions (Cont.)

- Result of aggregation does not have a name
 - ❖ Can use rename operation to give it a name
 - ❖ For convenience, we permit renaming as part of aggregate operation

branch-name $\text{g}_{\text{sum(balance) as sum-balance}}(\text{account})$





Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that do not match tuples in the other relation to the result of the join.
- Uses *null* values:
 - ❖ *null* signifies that the value is unknown or does not exist
 - ❖ All comparisons involving *null* are (roughly speaking) **false** by definition.
- Will study precise meaning of comparisons with nulls later



Outer Join – Example

- Relation *loan*

loan-number	branch-name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

- Relation *borrower*

customer-name	loan-number
Jones	L-170
Smith	L-230
Hayes	L-155





Outer Join – Example

■ Inner Join


$loan \bowtie Borrower$

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

■ Left Outer Join

$loan \leftarrow\bowtie Borrower$

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	null



Outer Join – Example

■ Right Outer Join


$loan \bowtie\rightarrow borrower$

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	null	null	Hayes

■ Full Outer Join

$loan \bowtie\bowtie borrower$


loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	null
L-155	null	null	Hayes





Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*.
- Aggregate functions simply ignore null values
 - ❖ Is an arbitrary decision. Could have returned null as result instead.
 - ❖ We follow the semantics of SQL in its handling of null values
- For duplicate elimination and grouping, null is treated like any other value, and two nulls are assumed to be the same
 - ❖ Alternative: assume each null is different from each other
 - ❖ Both are arbitrary decisions, so we simply follow SQL



Null Values


- Comparisons with null values return the special truth value *unknown*
 - ❖ If *false* was used instead of *unknown*, then $\text{not}(A < 5)$ would not be equivalent to $A \geq 5$
- Three-valued logic using the truth value *unknown*:
 - ❖ OR: $(\text{unknown} \text{ or } \text{true}) = \text{true}$,
 $(\text{unknown} \text{ or } \text{false}) = \text{unknown}$,
 $(\text{unknown} \text{ or } \text{unknown}) = \text{unknown}$
 - ❖ AND: $(\text{true} \text{ and } \text{unknown}) = \text{unknown}$,
 $(\text{false} \text{ and } \text{unknown}) = \text{false}$,
 $(\text{unknown} \text{ and } \text{unknown}) = \text{unknown}$
 - ❖ NOT: $(\text{not } \text{unknown}) = \text{unknown}$
 - ❖ In SQL “*P is unknown*” evaluates to true if predicate *P* evaluates to *unknown*
- Result of select predicate is treated as *false* if it evaluates to *unknown*





Modification of the Database

- The content of the database may be modified using the following operations:
 - ❖ Deletion
 - ❖ Insertion
 - ❖ Updating
- All these operations are expressed using the assignment operator.




Deletion

- A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database.
- Can delete only whole tuples; cannot delete values on only particular attributes
- A deletion is expressed in relational algebra by:

$$r \leftarrow r - E$$

where r is a relation and E is a relational algebra query.





Deletion Examples


- Delete all account records in the Perryridge branch.

$$account \leftarrow account - \sigma_{branch-name = "Perryridge"}(account)$$

- Delete all loan records with amount in the range of 0 to 50

$$loan \leftarrow loan - \sigma_{amount \geq 0 \text{ and } amount \leq 50}(loan)$$

- Delete all accounts at branches located in Needham.

$$\begin{aligned} r_1 &\leftarrow \sigma_{branch-city = "Needham"}(account \bowtie branch) \\ r_2 &\leftarrow \Pi_{branch-name, account-number, balance}(r_1) \\ r_3 &\leftarrow \Pi_{customer-name, account-number}(r_2 \bowtie depositor) \\ account &\leftarrow account - r_2 \\ depositor &\leftarrow depositor - r_3 \end{aligned}$$



Insertion

- To insert data into a relation, we either:
 - ❖ specify a tuple to be inserted
 - ❖ write a query whose result is a set of tuples to be inserted
- in relational algebra, an insertion is expressed by:

$$r \leftarrow r \cup E$$

where r is a relation and E is a relational algebra expression.

- The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple.






Insertion Examples

- Insert information in the database specifying that Smith has \$1200 in account A-973 at the Perryridge branch.

```
account ← account ∪ {("Perryridge", A-973, 1200)}  
depositor ← depositor ∪ {("Smith", A-973)}
```

- Provide as a gift for all loan customers in the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account.

```
r1 ← (σbranch-name = "Perryridge" (borrower ⋈ loan))  
account ← account ∪ Πbranch-name, account-number, 200 (r1)  
depositor ← depositor ∪ Πcustomer-name, loan-number (r1)
```




Updating

- A mechanism to change a value in a tuple without changing *all* values in the tuple
- Use the generalized projection operator to do this task

$$r \leftarrow \Pi_{F_1, F_2, \dots, F_i} (r)$$

- Each F_i is either
 - the i th attribute of r , if the i th attribute is not updated, or,
 - if the attribute is to be updated F_i is an expression, involving only constants and the attributes of r , which gives the new value for the attribute






Update Examples

- Make interest payments by increasing all balances by 5 percent.

$$\text{account} \leftarrow \Pi_{AN, BN, BAL * 1.05} (\text{account})$$

where *AN*, *BN* and *BAL* stand for *account-number*, *branch-name* and *balance*, respectively.

- Pay all accounts with balances over \$10,000 6 percent interest and pay all others 5 percent

$$\begin{aligned}\text{account} \leftarrow & \Pi_{AN, BN, BAL * 1.06} (\sigma_{BAL > 10000} (\text{account})) \\ & \cup \Pi_{AN, BN, BAL * 1.05} (\sigma_{BAL \leq 10000} (\text{account}))\end{aligned}$$


Views

- In some cases, it is not desirable for all users to see the entire logical model (i.e., all the actual relations stored in the database.)
- Consider a person who needs to know a customer's loan number but has no need to see the loan amount. This person should see a relation described, in the relational algebra, by

$$\Pi_{customer-name, loan-number} (borrower \times loan)$$

- Any relation that is not of the conceptual model but is made visible to a user as a "virtual relation" is called a **view**.






View Definition

- A view is defined using the **create view** statement which has the form

create view *v* as <query expression>

where <query expression> is any legal relational algebra query expression. The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
 - ❖ Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.



View Examples

- Consider the view (named *all-customer*) consisting of branches and their customers.

create view *all-customer* as

$$\begin{aligned} & \Pi_{branch-name, customer-name} (depositor \bowtie account) \\ & \cup \Pi_{branch-name, customer-name} (borrower \bowtie loan) \end{aligned}$$

- We can find all customers of the Perryridge branch by writing:

$$\begin{aligned} & \Pi_{customer-name} \\ & (\sigma_{branch-name = "Perryridge"} (all-customer)) \end{aligned}$$





Updates Through View


- Database modifications expressed as views must be translated to modifications of the actual relations in the database.
- Consider the person who needs to see all loan data in the *loan* relation except *amount*. The view given to the person, *branch-loan*, is defined as:

create view branch-loan as

$\Pi_{branch-name, loan-number} (loan)$


- Since we allow a view name to appear wherever a relation name is allowed, the person may write:

branch-loan \leftarrow *branch-loan* $\cup \{("Perryridge", L-37)\}$



Updates Through Views (Cont.)

- The previous insertion must be represented by an insertion into the actual relation *loan* from which the view *branch-loan* is constructed.
- An insertion into *loan* requires a value for *amount*. The insertion can be dealt with by either.
 - ⌚ rejecting the insertion and returning an error message to the user.
 - ⌚ inserting a tuple ("L-37", "Perryridge", *null*) into the *loan* relation
- Some updates through views are impossible to translate into database relation updates
 - ⌚ **create view v as** $\sigma_{branch-name = "Perryridge"} (account)$
 $v \leftarrow v \cup (L-99, Downtown, 23)$
- Others cannot be translated uniquely
 - ⌚ $all\text{-}customer} \leftarrow all\text{-}customer \cup \{("Perryridge", "John")\}$
⌚ Have to choose loan or account, and create a new loan/account number!





Views Defined Using Other Views

- One view may be used in the expression defining another view
- A view relation v_1 is said to *depend directly* on a view relation v_2 if v_2 is used in the expression defining v_1
- A view relation v_1 is said to *depend on* view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2
- A view relation v is said to be *recursive* if it depends on itself.



View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view v_1 be defined by an expression e_1 that may itself contain uses of view relations.
- View expansion of an expression repeats the following replacement step:

repeat

 Find any view relation v_i in e_1

 Replace the view relation v_i by the expression defining v_i

until no more view relations are present in e_1


- As long as the view definitions are not recursive, this loop will terminate





Tuple Relational Calculus

- A nonprocedural query language, where each query is of the form
$$\{t \mid P(t)\}$$
- It is the set of all tuples t such that predicate P is true for t
- t is a *tuple variable*, $t[A]$ denotes the value of tuple t on attribute A
- $t \in r$ denotes that tuple t is in relation r
- P is a *formula* similar to that of the predicate calculus




Predicate Calculus Formula

1. Set of attributes and constants
2. Set of comparison operators: (e.g., $<$, \leq , $=$, \neq , $>$, \geq)
3. Set of connectives: and (\wedge), or (\vee), not (\neg)
4. Implication (\Rightarrow): $x \Rightarrow y$, if x is true, then y is true

$$x \Rightarrow y \equiv \neg x \vee y$$

5. Set of quantifiers:


- $\exists t \in r (Q(t)) \equiv$ "there exists" a tuple in t in relation r such that predicate $Q(t)$ is true
- $\forall t \in r (Q(t)) \equiv Q$ is true "for all" tuples t in relation r





Banking Example

- *branch* (*branch-name*, *branch-city*, *assets*)
- *customer* (*customer-name*, *customer-street*, *customer-city*)
- *account* (*account-number*, *branch-name*, *balance*)
- *loan* (*loan-number*, *branch-name*, *amount*)
- *depositor* (*customer-name*, *account-number*)
- *borrower* (*customer-name*, *loan-number*)



Example Queries


- Find the *loan-number*, *branch-name*, and *amount* for loans of over \$1200

$$\{t \mid t \in \text{loan} \wedge t[\text{amount}] > 1200\}$$

- Find the loan number for each loan of an amount greater than \$1200

$$\{t \mid \exists s \in \text{loan} (t[\text{loan-number}] = s[\text{loan-number}] \wedge s[\text{amount}] > 1200)\}$$

Notice that a relation on schema [*loan-number*] is implicitly defined by the query






Example Queries

- Find the names of all customers having a loan, an account, or both at the bank

$$\{t \mid \exists s \in \text{borrower}(\text{t}[customer-name] = s[customer-name]) \\ \vee \exists u \in \text{depositor}(\text{t}[customer-name] = u[customer-name])\}$$

- Find the names of all customers who have a loan and an account at the bank


$$\{t \mid \exists s \in \text{borrower}(\text{t}[customer-name] = s[customer-name]) \\ \wedge \exists u \in \text{depositor}(\text{t}[customer-name] = u[customer-name])\}$$


Example Queries

- Find the names of all customers having a loan at the Perryridge branch

$$\{t \mid \exists s \in \text{borrower}(\text{t}[customer-name] = s[customer-name]) \\ \wedge \exists u \in \text{loan}(u[branch-name] = "Perryridge" \\ \wedge u[loan-number] = s[loan-number]))\}$$


- Find the names of all customers who have a loan at the Perryridge branch, but no account at any branch of the bank

$$\{t \mid \exists s \in \text{borrower}(\text{t}[customer-name] = s[customer-name]) \\ \wedge \exists u \in \text{loan}(u[branch-name] = "Perryridge" \\ \wedge u[loan-number] = s[loan-number])) \\ \wedge \text{not } \exists v \in \text{depositor}(v[customer-name] = \\ \text{t}[customer-name])\}$$





Example Queries

- Find the names of all customers having a loan from the Perryridge branch, and the cities they live in

$$\{t \mid \exists s \in \text{loan}(s[\text{branch-name}] = \text{"Perryridge"}) \wedge \exists u \in \text{borrower}(u[\text{loan-number}] = s[\text{loan-number}] \wedge t[\text{customer-name}] = u[\text{customer-name}]) \wedge \exists v \in \text{customer}(u[\text{customer-name}] = v[\text{customer-name}] \wedge t[\text{customer-city}] = v[\text{customer-city}]))\}$$


Example Queries

- Find the names of all customers who have an account at all branches located in Brooklyn:

$$\{t \mid \exists c \in \text{customer}(\text{f}[customer.name] = c[\text{customer-name}]) \wedge \forall s \in \text{branch}(s[\text{branch-city}] = \text{"Brooklyn"} \Rightarrow \exists u \in \text{account}(\text{s}[branch-name] = u[\text{branch-name}] \wedge \exists s \in \text{depositor}(\text{t}[customer-name] = s[\text{customer-name}] \wedge s[\text{account-number}] = u[\text{account-number}])))\}$$




Safety of Expressions

- It is possible to write tuple calculus expressions that generate infinite relations.
- For example, $\{t \mid \neg t \in r\}$ results in an infinite relation if the domain of any attribute of relation r is infinite
- To guard against the problem, we restrict the set of allowable expressions to safe expressions.
- An expression $\{t \mid P(t)\}$ in the tuple relational calculus is *safe* if every component of t appears in one of the relations, tuples, or constants that appear in P
 - ❖ NOTE: this is more than just a syntax condition.
 - ▀ E.g. $\{t \mid \#A=5 \vee \text{true}\}$ is not safe --- it defines an infinite set with attribute values that do not appear in any relation or tuples or constants in P .




Domain Relational Calculus

- A nonprocedural query language equivalent in power to the tuple relational calculus
- Each query is an expression of the form:

$$\{ < x_1, x_2, \dots, x_n > \mid P(x_1, x_2, \dots, x_n) \}$$

- ❖ x_1, x_2, \dots, x_n represent domain variables
- ❖ P represents a formula similar to that of the predicate calculus





Example Queries

- Find the *loan-number*, *branch-name*, and *amount* for loans of over \$1200

$$\{< l, b, a > \mid < l, b, a > \in \text{loan} \wedge a > 1200\}$$


- Find the names of all customers who have a loan of over \$1200

$$\{< c > \mid \exists l, b, a (< c, l > \in \text{borrower} \wedge < l, b, a > \in \text{loan} \wedge a > 1200)\}$$

- Find the names of all customers who have a loan from the Perryridge branch and the loan amount:

$$\{< c, a > \mid \exists l (< c, l > \in \text{borrower} \wedge \exists b (< l, b, a > \in \text{loan} \wedge b = \text{"Perryridge"}))\}$$

or $\{< c, a > \mid \exists l (< c, l > \in \text{borrower} \wedge < l, \text{"Perryridge"}, a > \in \text{loan})\}$




Example Queries

- Find the names of all customers having a loan, an account, or both at the Perryridge branch:

$$\begin{aligned} \{< c > \mid \exists l (&(< c, l > \in \text{borrower} \\ \wedge \exists b, a (< l, b, a > \in \text{loan} \wedge b = \text{"Perryridge"})) \\ \vee \exists a (< c, a > \in \text{depositor} \\ \wedge \exists b, n (< a, b, n > \in \text{account} \wedge b = \text{"Perryridge"}))\} \end{aligned}$$

- Find the names of all customers who have an account at all branches located in Brooklyn:

$$\begin{aligned} \{< c > \mid \exists s, n (&< c, s, n > \in \text{customer}) \wedge \\ \forall x, y, z (&< x, y, z > \in \text{branch} \wedge y = \text{"Brooklyn"}) \Rightarrow \\ \exists a, b (&< x, y, z > \in \text{account} \wedge < c, a > \in \text{depositor})\} \end{aligned}$$





Safety of Expressions

$$\{ < x_1, x_2, \dots, x_n > \mid P(x_1, x_2, \dots, x_n) \}$$

is safe if all of the following hold:

1. All values that appear in tuples of the expression are values from $\text{dom}(P)$ (that is, the values appear either in P or in a tuple of a relation mentioned in P).
2. For every “there exists” subformula of the form $\exists x (P_1(x))$, the subformula is true if and only if there is a value of x in $\text{dom}(P_1)$ such that $P_1(x)$ is true.
3. For every “for all” subformula of the form $\forall x (P_1(x))$, the subformula is true if and only if $P_1(x)$ is true for all values x from $\text{dom}(P_1)$.



End of Chapter 3




Result of $\sigma_{branch-name = "Perryridge"}(loan)$

loan-number	branch-name	amount
L-15	Perryridge	1500
L-16	Perryridge	1300



Loan Number and the Amount of the Loan

loan-number	amount
L-11	900
L-14	1500
L-15	1500
L-16	1300
L-17	1000
L-23	2000
L-93	500






Names of All Customers Who Have Either a Loan or an Account

customer-name


Adams
Curry
Hayes
Jackson
Jones
Smith
Williams
Lindsay
Johnson
Turner



Customers With An Account But No Loan

customer-name

Johnson
Lindsay
Turner





Result of $borrower \times loan$

customer-name	borrower. loan-number	loan. loan-number	branch-name	amount
Adams	L-16	L-11	Round Hill	900
Adams	L-16	L-14	Downtown	1500
Adams	L-16	L-15	Perryridge	1500
Adams	L-16	L-16	Perryridge	1300
Adams	L-16	L-17	Downtown	1000
Adams	L-16	L-23	Redwood	2000
Adams	L-16	L-93	Mianus	500
Curry	L-93	L-11	Round Hill	900
Curry	L-93	L-14	Downtown	1500
Curry	L-93	L-15	Perryridge	1500
Curry	L-93	L-16	Perryridge	1300
Curry	L-93	L-17	Downtown	1000
Curry	L-93	L-23	Redwood	2000
Curry	L-93	L-93	Mianus	500
Hayes	L-15	L-11		900
Hayes	L-15	L-14		1500
Hayes	L-15	L-15		1500
Hayes	L-15	L-16		1300
Hayes	L-15	L-17		1000
Hayes	L-15	L-23		2000
Hayes	L-15	L-93		500
...
...
...
Smith	L-23	L-11	Round Hill	900
Smith	L-23	L-14	Downtown	1500
Smith	L-23	L-15	Perryridge	1500
Smith	L-23	L-16	Perryridge	1300
Smith	L-23	L-17	Downtown	1000
Smith	L-23	L-23	Redwood	2000
Smith	L-23	L-93	Mianus	500
Williams	L-17	L-11	Round Hill	900
Williams	L-17	L-14	Downtown	1500
Williams	L-17	L-15	Perryridge	1500
Williams	L-17	L-16	Perryridge	1300
Williams	L-17	L-17	Downtown	1000
Williams	L-17	L-23	Redwood	2000
Williams	L-17	L-93	Mianus	500

©Silberschatz, Korth and Sudarshan

Database System Concepts

3.95



Result of $\sigma_{branch-name = "Perryridge"}(borrower \times loan)$

customer-name	borrower. loan-number	loan. loan-number	branch-name	amount
Adams	L-16	L-15	Perryridge	1500
Adams	L-16	L-16	Perryridge	1300
Curry	L-93	L-15	Perryridge	1500
Curry	L-93	L-16	Perryridge	1300
Hayes	L-15	L-15	Perryridge	1500
Hayes	L-15	L-16	Perryridge	1300
Jackson	L-14	L-15	Perryridge	1500
Jackson	L-14	L-16	Perryridge	1300
Jones	L-17	L-15	Perryridge	1500
Jones	L-17	L-16	Perryridge	1300
Smith	L-11	L-15	Perryridge	1500
Smith	L-11	L-16	Perryridge	1300
Smith	L-23	L-15	Perryridge	1500
Smith	L-23	L-16	Perryridge	1300
Williams	L-17	L-15	Perryridge	1500
Williams	L-17	L-16	Perryridge	1300

©Silberschatz, Korth and Sudarshan

Database System Concepts

3.96





Result of $\Pi_{customer-name}$

customer-name

Adams
Hayes



Result of the Subexpression

balance

500
400
700
750
350





Largest Account Balance in the Bank

<i>balance</i>
900



Customers Who Live on the Same Street and In the Same City as Smith

<i>customer-name</i>
Curry Smith






Customers With Both an Account and a Loan at the Bank


customer-name

Hayes
Jones
Smith



Result of $\Pi_{\text{customer-name, loan-number, amount}}$ $(\text{borrower} \bowtie \text{loan})$

<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Adams	L-16	1300
Curry	L-93	500
Hayes	L-15	1500
Jackson	L-14	1500
Jones	L-17	1000
Smith	L-23	2000
Smith	L-11	900
Williams	L-17	1000





Result of $\Pi_{branch-name}(\sigma_{customer-city = "Harrison"}(customer \bowtie account \bowtie depositor))$

<i>branch-name</i>
Brighton
Perryridge



Result of $\Pi_{branch-name}(\sigma_{branch-city = "Brooklyn"}(branch))$


<i>branch-name</i>
Brighton
Downtown






Result of $\Pi_{customer-name, branch-name} (depositor \bowtie account)$

customer-name	branch-name
Hayes	Perryridge
Johnson	Downtown
Johnson	Brighton
Jones	Brighton
Lindsay	Redwood
Smith	Mianus
Turner	Round Hill



The credit-info Relation


customer-name	branch-name
Hayes	Perryridge
Johnson	Downtown
Johnson	Brighton
Jones	Brighton
Lindsay	Redwood
Smith	Mianus
Turner	Round Hill






Result of $\Pi_{customer-name, (limit - credit-balance)} as credit-available (credit-info).$

customer-name	credit-available
Curry	250
Jones	5300
Smith	1600
Hayes	0



The *pt-works* Relation

employee-name	branch-name	salary
Adams	Perryridge	1500
Brown	Perryridge	1300
Gopal	Perryridge	5300
Johnson	Downtown	1500
Loreena	Downtown	1300
Peterson	Downtown	2500
Rao	Austin	1500
Sato	Austin	1600






The *pt-works* Relation After Grouping

<i>employee-name</i>	<i>branch-name</i>	<i>salary</i>
Rao	Austin	1500
Sato	Austin	1600
Johnson	Downtown	1500
Loreena	Downtown	1300
Peterson	Downtown	2500
Adams	Perryridge	1500
Brown	Perryridge	1300
Gopal	Perryridge	5300



Result of $\text{branch-name} \Sigma \text{sum}(\text{salary})$ (*pt-works*)

<i>branch-name</i>	<i>sum of salary</i>
Austin	3100
Downtown	5300
Perryridge	8100





Result of $\text{branch-name} \sum \text{sum salary}, \max(\text{salary}) \text{ as max-salary}$ (*pt-works*)


<i>branch-name</i>	<i>sum-salary</i>	<i>max-salary</i>
Austin	3100	1600
Downtown	5300	2500
Perryridge	8100	5300



The employee and *ft-works* Relations

<i>employee-name</i>	<i>street</i>	<i>city</i>
Coyote	Toon	Hollywood
Rabbit	Tunnel	Carrotville
Smith	Revolver	Death Valley
Williams	Seaview	Seattle


<i>employee-name</i>	<i>branch-name</i>	<i>salary</i>
Coyote	Mesa	1500
Rabbit	Mesa	1300
Gates	Redmond	5300
Williams	Redmond	1500





The Result of $employee \bowtie ft\text{-}works$

employee-name	street	city	branch-name	salary
Coyote	Toon	Hollywood	Mesa	1500
Rabbit	Tunnel	Carrotville	Mesa	1300
Williams	Seaview	Seattle	Redmond	1500



The Result of $employee \bowtie ft\text{-}works$

employee-name	street	city	branch-name	salary
Coyote	Toon	Hollywood	Mesa	1500
Rabbit	Tunnel	Carrotville	Mesa	1300
Williams	Seaview	Seattle	Redmond	1500
Smith	Revolver	Death Valley	null	null






Result of employee \bowtie ft-works

employee-name	street	city	branch-name	salary
Coyote	Toon	Hollywood	Mesa	1500
Rabbit	Tunnel	Carrotville	Mesa	1300
Williams	Seaview	Seattle	Redmond	1500
Gates	null	null	Redmond	5300



Result of employee \bowtie ft-works

employee-name	street	city	branch-name	salary
Coyote	Toon	Hollywood	Mesa	1500
Rabbit	Tunnel	Carrotville	Mesa	1300
Williams	Seaview	Seattle	Redmond	1500
Smith	Revolver	Death Valley	null	null
Gates	null	null	Redmond	5300





Tuples Inserted Into *loan* and *borrower*


<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-11	Round Hill	900
L-14	Downtown	1500
L-15	Perryridge	1500
L-16	Perryridge	1300
L-17	Downtown	1000
L-23	Redwood	2000
L-93	Mianus	500
<i>null</i>	<i>null</i>	1900

<i>customer-name</i>	<i>loan-number</i>
Adams	L-16
Curry	L-93
Hayes	L-15
Jackson	L-14
Jones	L-17
Smith	L-11
Smith	L-23
Williams	L-17
Johnson	<i>null</i>

Database System Concepts

3.117

©Silberschatz, Korth and Sudarshan




Names of All Customers Who Have a Loan at the Perryridge Branch

<i>customer-name</i>
Adams
Hayes


Database System Concepts

3.118

©Silberschatz, Korth and Sudarshan



E-R Diagram



Database System Concepts

3.119

©Silberschatz, Korth and Sudarshan

The *branch* Relation

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>
Brighton	Brooklyn	7100000
Downtown	Brooklyn	9000000
Mianus	Horseneck	400000
North Town	Rye	3700000
Perryridge	Horseneck	1700000
Pownal	Bennington	300000
Redwood	Palo Alto	2100000
Round Hill	Horseneck	8000000

Database System Concepts


3.120

©Silberschatz, Korth and Sudarshan



The *loan* Relation

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-11	Round Hill	900
L-14	Downtown	1500
L-15	Perryridge	1500
L-16	Perryridge	1300
L-17	Downtown	1000
L-23	Redwood	2000
L-93	Mianus	500



The *borrower* Relation


<i>customer-name</i>	<i>loan-number</i>
Adams	L-16
Curry	L-93
Hayes	L-15
Jackson	L-14
Jones	L-17
Smith	L-11
Smith	L-23
Williams	L-17





Chapter 7: SQL

- Basic Structure
- Simple Queries
- Nested Subqueries
- Aggregate Functions
- Set Operations
- With Clause
- Views
- Modification of the Database
- Joined Relations
- Data Definition Language
- Embedded SQL, ODBC and JDBC





Basic Structure

- SQL is based on set and relational operations with certain modifications and enhancements
- A typical SQL query has the form:

```
select A1, A2, ..., An
  from r1, r2, ..., rm
  where predicate
```

 - ❖ A_is represent attributes
 - ❖ r_is represent relations (tables)
 - ❖ predicate is any predicate.
- This query is equivalent to the relational algebra expression.
$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$
- The result of an SQL query is a relation.
- **NOTE:** SQL does not permit the '-' character in names. SQL names are case insensitive, i.e. you can use capital or small letters.





Simple Queries (2)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get supplier numbers and status for suppliers in Paris in descending order of status.

```
select S#, Status
from S
where City = 'Paris'
order by Status desc ;
```

Result:

S#	Status
S3	30
S2	10

For all blue parts, get the weights in grams.

```
select P#, Weight*454
from P
where Color = 'Blue'
order by 2, P# ;
```

Result:


P#	Weight
P5	5448
P3	7718

Include constant in select clause.

```
select P#, 'Weights in grams = ', Weight*454
from P
where Color = 'Blue' ;
```

Result:

```
P#
P3    Weights in grams = 7718
P5    Weights in grams = 5448
```



Simple Queries (between)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get parts whose weight is in range 16 to 19 (inclusive).

```
select *
from P
where Weight between 16 and 19 ;
```

Result:

P#	Pname	Color	Weight	City
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P6	Cog	Red	19	London

Get parts whose weight is not in range 16 to 19.

```
select P#, Pname, Color, Weight, City
from P
where Weight not between 16 and 19 ;
```

Result:

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris



Simple Queries (in)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get parts whose weight is in range 16 to 19 (inclusive).

```
select *
from P
where Weight in {12, 16, 17} ;
```

Result:


P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P5	Cam	Blue	12	Paris

Get parts whose weight is not in range 16 to 19.

```
select P#, Pname, Color, Weight, City
from P
where Weight not in {12, 16, 17} ;
```

Result:

P#	Pname	Color	Weight	City
P4	Screw	Red	14	London
P6	Cog	Red	19	London



Simple Queries (like)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get parts whose names begin with the letter C.

```
select *
from P
where Pname like 'C*' ;
```

Result:

P#	Pname	Color	Weight	City
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

% stands for any string, ? stands for any character

Sname like '?la*' – all supplier names with second character l and third character a.

Pname like '????' – all part names 4 character long.

City not like '*o*' – all city names which does not contain character o.

like 'Main*' escape '\' – match Main*

SQL supports a variety of string operations such as: concatenation ("||"), converting from upper to lower case (and vice versa), finding string length, extracting substrings, etc.





Simple Queries (null values)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get parts whose color is not null.

```
select *
from P
where Color is not null ;
```

Result:

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

null signifies an unknown value or that a value does not exist.

The result of any arithmetic expression involving null is null (E.g. 5 + null returns null).

Any comparison with null returns unknown (E.g. 5 < null or null <> null or null = null).

©Silberschatz, Korth and Sudarshan



Simple Queries (natural join)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get all combination suppliers - parts located in the same city.

```
select S.* , P.*
from S, P
where S.City = P.City ;
```

Result:

S#	Sname	Status	S.City	P#	Pname	Color	Weight	P.City
S1	Smith	20	London	P1	Nut	Red	12	London
S1	Smith	20	London	P4	Screw	Red	14	London
S1	Smith	20	London	P6	Cog	Red	19	London
S2	Jones	10	Paris	P2	Bolt	Green	17	Paris
S2	Jones	10	Paris	P5	Cam	Blue	12	Paris
S3	Blake	30	Paris	P2	Bolt	Green	17	Paris
S3	Blake	30	Paris	P5	Cam	Blue	12	Paris
S4	Clark	20	London	P5	Nut	Red	12	London
S4	Clark	20	London	P5	Screw	Red	14	London
S4	Clark	20	London	P5	Cog	Red	19	London

How conceptually join is constructed:

- Form the cartesian product of the tables listed in **from** clause (in our example the new table will have $5 \cdot 6 = 30$ rows)
- Eliminate from the cartesian product all those rows that do not satisfy join predicate (**where** clause)



©Silberschatz, Korth and Sudarshan

Simple Queries (natural join)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	300	S2	P1
S1	P2	200	S2	P2
S1	P3	400	S3	P2
S1	P4	200	S4	P2
S1	P5	100	S4	P4
S1	P6	100	S4	P5
				400

Same, but supplier city follows part city (alphabetically).

```
select S.* , P.*  
from S, P  
where S.City > P.City ;
```

Result:

S#	Sname	Status	S.City	P#	Pname	Color	Weight	P.City
S2	Jones	10	Paris	P1	Nut	Red	12	London
S2	Jones	10	Paris	P4	Screw	Red	14	London
S2	Jones	10	Paris	P6	Cog	Red	19	London
S3	Blake	30	Paris	P1	Nut	Red	12	London
S3	Blake	30	Paris	P4	Screw	Red	14	London
S3	Blake	30	Paris	P6	Cog	Red	19	London

Get all combination suppliers - parts located in the same city, without suppliers that have status 20.

```
select S.* , P.*  
from S, P  
where S.City = P.City and S.Status <> 20 ;
```

Result:

S#	Sname	Status	S.City	P#	Pname	Color	Weight	P.City
S2	Jones	10	Paris	P2	Bolt	Green	17	Paris
S2	Jones	10	Paris	P5	Cam	Blue	12	Paris
S3	Blake	30	Paris	P2	Bolt	Green	17	Paris
S3	Blake	30	Paris	P5	Cam	Blue	12	Paris



Simple Queries (natural join)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	300	S2	P1
S1	P2	200	S2	P2
S1	P3	400	S3	P2
S1	P4	200	S4	P2
S1	P5	100	S4	P4
S1	P6	100	S4	P5
				400

Get all pairs of city names such that a supplier located in the first city supplies a part stored in the second city.


For example, supplier S1 supplies part P1; suppliers S1 is located in London, and part P1 is stored in London; so 'London, London' is a pair of cities in the result.

```
select distinct S.City, P.City  
from S, SP, P  
where S.S# = SP.S# and SP.P# = P.P# ;
```

Result:

S.City	P.City
London	London
London	Paris
London	Rome
Paris	London
Paris	Paris

This example shows join of 3 tables.



Simple Queries (join a table with itself)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get all pairs of supplier numbers such that the two suppliers are co-located.

```
select Sup1.S#, Sup2.S#
from S as Sup1, S as Sup2
where Sup1.City = Sup2.City ;
```

Result:


S#	S#
S1	S1
S1	S4
S2	S2
S2	S3
S3	S2

This result can be cleared up as follows:

```
select Sup1.S#, Sup2.S#
from S as Sup1, S as Sup2
where Sup1.City = Sup2.City and Sup1.S# > Sup2.S# ;
```

Result:

S#	S#
S1	S4
S2	S3



SubQueries

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

Get suppliers names for suppliers who supplies part P2.

```
select S.Sname
from S
where S.S# in ( select SP.S#
from SP
where SP.P# = 'P2' ) ;
```

Result:

Sname
Smith
Jones
Blake
Clark

The nested subqueries are evaluated first.

So, our query is equivalent to:

```
select S.Sname
from S
where S.S# in ('S1', 'S2', 'S3', 'S4');
```

The same using join.

```
select S.Sname
from S, SP
where S.S# = SP.S# and SP.P# = 'P2' ;
```

The join of S and SP over supplier numbers is a table of 12 rows from which we select those 4 rows that have the part number P2.



SubQueries (correlated)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get suppliers names for suppliers who supplies part P2.

```
select Sname
from S
where P2' in ( select P#
from SP
where S#= S.S# );
```


Result:

Sname
Smith
Jones
Blake
Clark

In the last line the unqualified reference *S#* is implicitly qualified by *SP*. Here, inner subquery cannot be evaluated once and for all before the outer query is evaluated (variable *S.S#* is unknown). Such subqueries are called *correlated*. The system examines one by one rows of table *S* and each time evaluate the subquery.

Some people prefer to use aliases in correlated subqueries.

```
select SX.Sname
from S as SX
where P2' in ( select P#
from SP
where S#= SX.S# );
```



SubQueries (more nesting)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get suppliers names for suppliers who supply at least one red part.

```
select Sname
from S
where S# in ( select S#
from SP
where P# in ( select P#
from P
where Color = 'Red' ) );
```

Result:
Sname
Smith
Jones
Clark

The innermost subquery evaluates to the set {‘P1’, ‘P4’, ‘P6’}. The next subquery evaluates in turn to the set {‘S1’, ‘S2’, ‘S4’}. Last, the outermost select evaluates to the final result. In general, subqueries can be nested to any depth.

The same using join.

```
select distinct S.Sname
from S, SP, P
where S.S#= SP.S# and SP.P#= P.P#
and P.Color = 'Red' ;
```



SubQueries (with same table)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get supplier numbers for suppliers who supply at least one part supplied by supplier S2.

```
select distinct S#
from SP
where P# in ( select P#
               from SP
               where S# = 'S2' );
```

Result:

S#
S1
S2
S3
S4

The reference *SP* in the subquery does not mean the same thing as reference to *SP* in the outer query. They are different variables. Using aliases will make this fact explicit.

The same using join.

```
select distinct SP1.S#
from SP as SP1, SP as SP2
where SP1.P# = SP2.P#
      and SP2.S# = 'S2';
```



SubQueries (correlated with same table)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get part numbers for all parts supplied by more than one supplier.

```
select distinct SP1.P#
from SP as SP1
where SP1.P# in ( select SP2.P#
                     from SP as SP2
                     where SP2.S# = SP1.S# );
```

Result:

P#
P1
P2
P4
P5


Operation of this query: For each row in turn, *SP1* of table *SP*, extract the *P#* value, iff that *P#* value appears in some row *SP2* of table *SP* whose *S#* value is not that in row *SP1*. Note that at least one alias must be used, but not both.

Get supplier numbers for suppliers who are located in the same city as supplier S1.

```
select S#
from S
where City = ( select City
                  from S
                  where S# = 'S1' );
```

Result:

S#
S1
S4



SubQueries (exists)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	300	S2	P1 300
S1	P2	200	S2	P2 400
S1	P3	400	S3	P2 200
S1	P4	200	S4	P2 200
S1	P5	100	S4	P4 300
S1	P6	100	S4	P5 400

Get suppliers names for suppliers who supplies part P2.

```
select Sname
from S
where exists ( select *
               from SP
               where S# = S.S# and P# = 'P2' );
```

Result:

Sname

Smith

Jones

Blake

Clark


Predicate **exists x (predicate-involving-x)** is true iff **predicate-involving-x** is true for some x. For example if x=1,2,...,10 then **exists x (x<5)** is true, while **exists x (x<0)** is false.

Get suppliers names for suppliers who do not supply part P2.

```
select Sname
from S
```

where not exists (select *

In general, **exists** is one of the most important SQL feature. In fact, any query expressed using **in** can be formulated using **exists**. The converse is not true.



from SP

where S# = S.S# and P# = 'P2');

Result:

Sname

Adams

©Silberschatz, Korth and Sudarshan



SubQueries (not exists)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	300	S2	P1 300
S1	P2	200	S2	P2 400
S1	P3	400	S3	P2 200
S1	P4	200	S4	P2 200
S1	P5	100	S4	P4 300
S1	P6	100	S4	P5 400

Get supplier names for suppliers who supply all parts.

```
select Sname
from S
where not exists
( select *
   from P
   where not exists
( select *
   from SP
   where S# = S.S# and P# = p.pp );
```

Result:

Sname

Smith

The query can be paraphrased according to the above SQL statement: *Select supplier names for supplier such that there does not exist a part that they do not supply.*



SubQueries (all, some)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get the all part numbers that have greater shipment quantity than all parts located in London.

```
select P#
from SP
where QTY > all
      ( select QTY
        from SP, P
        where City = 'London' ) ;
```

P#
P3
P2
P5

Get the all part numbers that have greater shipment quantity than some part located in London. Result:

```
select P#
from SP
where QTY > some
      ( select QTY
        from SP, P
        where City = 'London' ) ;
```

P#
P1
P2
P3
P4
P5



Definition of Some and All Clauses

$$(5 < \text{some} \begin{array}{|c|}\hline 0 \\ \hline 5 \\ \hline 6 \end{array}) = \text{true} \quad (5 < \text{some} \begin{array}{|c|}\hline 0 \\ \hline 5 \end{array}) = \text{false}$$


$$(5 = \text{some} \begin{array}{|c|}\hline 0 \\ \hline 5 \end{array}) = \text{true} \quad (5 \neq \text{some} \begin{array}{|c|}\hline 0 \\ \hline 5 \end{array}) = \text{true} \text{ (since } 0 \neq 5\text{)}$$

$(= \text{some}) \equiv \text{in}$. However, $(\neq \text{some}) \not\equiv \text{not in}$

$$(5 < \text{all} \begin{array}{|c|}\hline 0 \\ \hline 5 \\ \hline 6 \end{array}) = \text{false} \quad (5 < \text{all} \begin{array}{|c|}\hline 6 \\ \hline 10 \end{array}) = \text{true}$$

$$(5 = \text{all} \begin{array}{|c|}\hline 4 \\ \hline 5 \end{array}) = \text{false} \quad (5 \neq \text{all} \begin{array}{|c|}\hline 4 \\ \hline 6 \end{array}) = \text{true} \text{ (since } 5 \neq 4 \text{ and } 5 \neq 6\text{)}$$

$(\neq \text{all}) \equiv \text{not in}$. However, $(= \text{all}) \not\equiv \text{in}$



Aggregate Functions (count, sum, max)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get the number of shipments for part P2.

```
select count(*)          Result:  
from SP  
where P# = 'P2' ;        4
```

Get the total quantity of part P2 supplied.

```
select sum(QTY)           Result:  
from SP  
where P# = 'P2' ;        1000
```

Get supplier numbers for suppliers with status less than current maximum status.

```
select S#                  Result:  
from S  
where Status <           S#  
                           ( select max(Status)  
                             from S ) ;      S1  
                                         S2  
                                         S4
```



Aggregate Functions (min, avg)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens


P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

Get the all part names for parts with minimum weights.

```
select Pname                Result:  
from P  
where Weight =             Pname  
                           ( select min(Weight)  
                             from P ) ;      Nut  
                                         Cam
```

Get supplier numbers, status nad city for all suppliers whose status is greater than or equal to the average for their city.

```
select S#, Status, City      Result:  
from S as S1  
where Status >=            S#  Status  City  
                           ( select avg(Status)  
                             from S as S2  
                           where S2.City = S1.City ) ;  S1 20  London  
                                         S3 30  Paris  
                                         S4 30  London  
                                         S5 30  Athens
```





Aggregate Functions (group by)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get the total quantity supplied for each part.

```
select P#, sum(QTY)
from SP
group by P#;
```

Result:

P#
P1
P2
P3
P4
P5
P6

©Silberschatz, Korth and Sudarshan

For each part supplied, get the part number and the total quantity supplied of that part, excluding shipment from supplier S1.

```
select P#, sum(QTY)
from SP
where S# <> 'S1'
group by P#;
```

Result:

P#
P1
P2
P4
P5



Aggregate Functions (having)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get part numbers for all parts supplied by more than one supplier.

```
select P#
from SP
group by P#
having count(*) > 1 ;
```


Result:

P#
P1
P2
P4
P5

Having is to groups what **where** is to rows. (If **having** is specified, **group by** should be also specified). **Having** is used to eliminate groups just as **where** is used to eliminate rows.

The same without **group by/having**.

```
select P#,
from P
where 1 < ( select count(S#)
from SP
where P# = P.P# );
```



Set Operations (union)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Get part numbers for parts with weight more than 16 pounds or are supplied by supplier S2.

Result:
select P#
from P
where Weight > 16 union select P#
from SP
where S# = 'S2' ;

Since a relation is set of rows, it is possible to construct union, intersection and difference between them. However, to be result a relation the two original relation must be set-compatable:

1. *to have the same number of columns.*
2. *the i-th column of both relations must have the same data type.*

The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations \cup , \cap , $-$.

Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.



Set Operations (intersect, except)

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400


Get supplier numbers for suppliers who supply part P1 and are located in London.

Result:
select S#
from SP
where P# = 'P1' intersect select S#
from S
where City = 'London' ;

Get supplier numbers for suppliers who supply part P2 and are not located in London.

select S#
from SP
where P# = 'P2' except select S#
from S
where City = 'London' ;

Result:
S#
S2
S3



A Comprehensive Example

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	300	S2	P1 300
S1	P2	200	S2	P2 400
S1	P3	400	S3	P2 200
S1	P4	200	S4	P2 200
S1	P5	100	S4	P4 300
S1	P6	100	S4	P5 400

For all red and blue parts such that the total quantity supplied is greater than 350 (excluding from the total all shipments for which the quantity is less than or equal to 200), get the part number, the weight in grams, the color, and the maximum supplied of that part. Order the result by decreasing part number within ascending values of that maximum.

```
select P.P#, Weight in grams = ', P.Weight*454,
       P.Color, MSQuantity = ', max (SP.QTY)
from P, SP
where P.P# = SP.P#
      and P.Color in ('Red', 'Blue')
      and SP.QTY > 200
group by P.P#, P.Weight, P.Color
having sum (QTY) > 350
order by 6, P.P#, desc ;
```

Result:

P#	Color	MSQuantity
P1	Red	MSQuantity = 300
P5	Blue	MSQuantity = 400
P3	Blue	MSQuantity = 400



With Clause

Get all supplier names with maximum status.

```
with maxst(value) as
  select max(Status)
  from S
select Sname
from S
where Status = maxst.value;
```

Result:

Sname
Blake
Adams

With clause allows views to be defined locally to a query, rather than globally. Analogous to procedures in a programming language.

Get all part numbers where the total their shipments is greater than the average of the total supplier shipments at all suppliers.

```
with ptotal(P#, value) as
  select P#, sum(QTY)
  from SP
  group by P#
with pavg(S#, value) as
  select S#, avg(QTY)
  from SP
  group by P#
select P#
from ptotal, pavg
where ptotal.value > pavg.value;
```



Derived Relations

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400


Get the average quantity of those supplier shipments where the average quantity is greater than 250.

```
select S#, AvgShip
  from (select S#, avg (QTY)
        from SP
       group by S#)
    as result (S#, AvgShip)
   where AvgShip > 250
```

Result:

S#	AvgShip
S2	350
S4	300

Note that we do not need to use the **having** clause, since we compute the temporary (view) relation *result* in the **from** clause, and the attributes of *result* can be used directly in the **where** clause.



Views

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Create view from good suppliers (with status greater than 15).

```
create view GoodSup
  as select S#, Status,City
    from S
   where Status > 15 ;
```

S#	Status	City
S1	20	London
S3	30	Paris
S4	20	London
S5	30	Athens

GoodSup is in effect a “window” into real table S. The window is dynamic because changes of S is automatically visible through the window GoodSup. Some users may genuinely believe that GoodSup is a “real” table.

Provide a mechanism to hide certain data from the view of certain users. To create a view we use the command:

create view v as <query expression>

where:

- <query expression> is any legal expression
- the view name is represented by v



Views

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Query on view (suppliers not located in London).

Result:

```
select S#, City
from GoodSup
where City <> 'London' ;
```

S#	City
S3	Paris
S5	Athens

Create view of part numbers and names for parts with weight more than 16 pounds or are supplied by supplier S2.

```
select P#, Pname
from P
where Weight > 16 union
      select distinct P#, Pname
from P, SP
where P.P# = SP.P#
and S# = 'S2' ;
```

Result:

P#	Pname
P1	Nut
P2	Bolt
P3	Screw
P6	Cog



Modification of the Database – Deletion

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

Delete all suppliers in Paris.

```
delete S#, City
from S
where City = 'Paris' ;
```

Delete all shipments.

```
delete
from SP ;
```

Delete all shipments for suppliers in London.

```
delete
from SP
where 'London' = ( select City
from S
where S.S# = SP.S# )
```



Modification of the Database – Deletion

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

General form of **delete** statement:

```
delete
from table
[ where predicate ]
```

Delete all shipments with quantity below the average.

```
delete
from SP
where QTY < ( select avg(QTY)
    from SP );
```

Problem: **as we delete tuples from SP, the average quantity changes**

Solution used in SQL:

1. First, compute **avg** balance and find all tuples to delete
2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)



Modification of the Database – Insertion

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Add part P7 with unknown name and color.

```
insert
into P (P#, City, Weight)
values ('P7', 'Athens', 2);
```

Name and color will have null values.

Add part P8 to table P.

```
insert
into P
values ('P8', 'Sprocket', 'Pink', 14, 'Nice');
```

Add a new shipment with supplier S20, part number P20, and quantity 1000.

```
insert
into SP (S#, P#, QTY)
values ('S20', 'P20', 1000);
```



Modification of the Database – Insertion

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	300
S1	P2	200	S2	P2	400
S1	P3	400	S3	P2	200
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400


General form of **insert** statement:

```
insert
into table [ (field1, field2, field3, ...) ]
values ( constant1, constant2, constant3, ... ) ;      or
insert
into table [ (field1, field2, field3, ...) ]
subquery ;
```

For each part supplied, get the part number and the total quantity, and save the result in the database.

```
create table temp
( P#    char(6)
  TOTQTY integer ) ;

insert into temp ( P#, TOTQTY )
select P#, sum(QTY)
from SP
group by P#;
```



Modification of the Database – Updates

S#	Sname	Status	City
S1	Smith	40	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	40	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Yellow	22	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY			
S1	P1	300	S2	P1	0
S1	P2	200	S2	P2	0
S1	P3	400	S3	P2	0
S1	P4	200	S4	P2	200
S1	P5	100	S4	P4	300
S1	P6	100	S4	P5	400

Double status for all suppliers in London.


```
update S
set status = status * 2
where city = 'London' ;
```

Change the color and weight of part P2.

```
update P
set color = 'Yellow', weight = weight + 5
where P# = 'P2' ;
```

Set the shipment quantity to zero for all suppliers in Paris.

```
update SP
set QTY = 0
where 'Paris' = ( select city
                  from S
                 where S.S# = SP.S# ) ;
```



Modification of the Database – Updates

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	318	S2	P1 318
S1	P2	210	S2	P2 424
S1	P3	424	S3	P2 210
S1	P4	210	S4	P2 210
S1	P5	105	S4	P4 318
S1	P6	105	S4	P5 424

General form of **update** statement:

```
update table
set field = expression
[ , field = expression ] ...
where predicate ;
```

Increase all shipment quantities over 200 by 6%, and all others by 5%.


```
update SP
set QTY = QTY * 1.06
where QTY > 200
```

```
update account
set QTY = QTY * 1.05
where QTY <= 200
```

The order is important

Can be done better using the **case** statement

```
update SP
set QTY = case
when QTY <= 200 then QTY * 1.05
else QTY * 1.06
end
```



Modification of the Views

S#	Sname	Status	City
S1	Smith	20	London
S2	Jones	10	Paris
S3	Blake	30	Paris
S4	Clark	20	London
S5	Adams	30	Athens

P#	Pname	Color	Weight	City
P1	Nut	Red	12	London
P2	Bolt	Green	17	Paris
P3	Screw	Blue	17	Rome
P4	Screw	Red	14	London
P5	Cam	Blue	12	Paris
P6	Cog	Red	19	London

S#	P#	QTY		
S1	P1	300	S2	P1 300
S1	P2	200	S2	P2 400
S1	P3	400	S3	P2 200
S1	P4	200	S4	P2 200
S1	P5	100	S4	P4 300
S1	P6	100	S4	P5 400

Create a view of shipment relation (SP), hiding the QTY attribute.

```
create view Ship as
select S#, P#
from SP
```

Add a new shipment to ship.

```
insert
into Ship
values ('S5', 'P6');
```

- Updates on more complex views are difficult or impossible to translate, and hence are disallowed.
- Most SQL implementations allow updates only on simple views (without aggregates) defined on a single relation





Transactions

- A transaction is a sequence of queries and update statements executed as a single unit
 - ❖ Transactions are started implicitly and terminated by one of
 - **commit work:** makes all updates of the transaction permanent in the database
 - **rollback work:** undoes all updates performed by the transaction.
- Motivating example
 - ❖ Transfer of money from one account to another involves two steps:
 - deduct from one account and credit to another
 - ❖ If one step succeeds and the other fails, database is in an inconsistent state
 - ❖ Therefore, either both steps should succeed or neither should
- If any step of a transaction fails, all work done by the transaction can be undone by **rollback work**.
- Rollback of incomplete transactions is done automatically, in case of system failures



Transactions (Cont.)

- In most database systems, each SQL statement that executes successfully is automatically committed.
 - ❖ Each transaction would then consist of only a single statement
 - ❖ Automatic commit can usually be turned off, allowing multi-statement transactions, but how to do so depends on the database system
 - ❖ Another option in SQL:1999: enclose statements within

```
begin atomic  
...  
end
```




Joined Relations

- Join operations take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- Join condition – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- Join type – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

Join Types
inner join
left outer join
right outer join
full outer join

Join Conditions
natural
on <predicate>
using (A_1, A_2, \dots, A_n)



Joined Relations – Datasets for Examples


- Relation *loan*

loan-number	branch-name	amount
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

- Relation *borrower*

customer-name	loan-number
Jones	L-170
Smith	L-230
Hayes	L-155

- Note: borrower information missing for L-260 and loan information missing for L-155






Joined Relations – Examples

- *loan inner join borrower on
loan.loan-number = borrower.loan-number*

loan-number	branch-name	amount	customer-name	loan-number
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230

- *loan left outer join borrower on
loan.loan-number = borrower.loan-number*

loan-number	branch-name	amount	customer-name	loan-number
L-170	Downtown	3000	Jones	L-170
L-230	Redwood	4000	Smith	L-230
L-260	Perryridge	1700	null	null




Joined Relations – Examples

- *loan natural inner join borrower*

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

- *loan natural right outer join borrower*

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	null	null	Hayes





Joined Relations – Examples

- *loan full outer join borrower using (loan-number)*

loan-number	branch-name	amount	customer-name
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	null	null	Hayes

- Find all customers who have either an account or a loan (but not both) at the bank.

```
select customer-name  
      from (depositor natural full outer join borrower)  
            where account-number is null or loan-number is null
```



Data Definition Language (DDL)

Allows the specification of not only a set of relations but also information about each relation, including:

- The schema for each relation.
- The domain of values associated with each attribute.
- Integrity constraints
- The set of indices to be maintained for each relations.
- Security and authorization information for each relation.
- The physical storage structure of each relation on disk.






Domain Types in SQL


- **char(n).** Fixed length character string, with user-specified length n .
- **varchar(n).** Variable length character strings, with user-specified maximum length n .
- **int.** Integer (a finite subset of the integers that is machine-dependent).
- **smallint.** Small integer (a machine-dependent subset of the integer domain type).
- **numeric(p,d).** Fixed point number, with user-specified precision of p digits, with n digits to the right of decimal point.
- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(n).** Floating point number, with user-specified precision of at least n digits.
- Null values are allowed in all the domain types. Declaring an attribute to be **not null** prohibits null values for that attribute.
- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person-name char(20) not null
```



Date/Time Types in SQL (Cont.)

- **date.** Dates, containing a (4 digit) year, month and date
 - ⌚ E.g. `date '2001-7-27'`
- **time.** Time of day, in hours, minutes and seconds.
 - ⌚ E.g. `time '09:00:30'` `time '09:00:30.75'`
- **timestamp:** date plus time of day
 - ⌚ E.g. `timestamp '2001-7-27 09:00:30.75'`
- **Interval:** period of time
 - ⌚ E.g. `Interval '1' day`
 - ⌚ Subtracting a date/time/timestamp value from another gives an interval value
 - ⌚ Interval values can be added to date/time/timestamp values
- Can extract values of individual fields from date/time/timestamp
 - ⌚ E.g. `extract (year from rstarttime)`
- Can cast string types to date/time/timestamp
 - ⌚ E.g. `cast <string-valued-expression> as date`





Create Table Construct


- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,  
    (integrity-constraint1),  
    ...  
    (integrity-constraintk))
```

- r is the name of the relation
- each A_i is an attribute name in the schema of relation r
- D_i is the data type of values in the domain of attribute A_i

- Example:

```
create table S  
(S#      char(5) not null,  
 Sname   char(20),  
 Status  smallint,  
 City    char(15) );
```




Integrity Constraints in Create Table

- not null**
- primary key** (A_1, \dots, A_n)
- check** (*Predicate*), where *Predicate* is a predicate

Example: Declare table P (Parts).

```
create table P  
( P#      char(6) not null,  
 Pname   char(20),  
 Color    char(10),  
 Weight   smallint,  
 City    char(15),  
 primary key (P#),  
 check (Weight >= 0))
```



primary key declaration on an attribute automatically ensures **not null** in SQL-92 onwards, needs to be explicitly stated in SQL-89



Drop and Alter Table Constructs

- The **drop table** command deletes all information about the dropped relation from the database.
- The **alter table** command is used to add attributes to an existing relation.

alter table r add $A D$

where A is the name of the attribute to be added to relation r and D is the domain of A .


☞ All tuples in the relation are assigned *null* as the value for the new attribute.

- The **alter table** command can also be used to drop attributes of a relation

alter table r drop A

where A is the name of an attribute of relation r

☞ Dropping of attributes not supported by many databases




Chapter 5: Relational Database Design



Chapter 5: Relational Database Design

- First Normal Form
- Pitfalls in Relational Database Design
- Functional Dependencies
- Decomposition
- Boyce-Codd Normal Form
- Third Normal Form
- Multivalued Dependencies and Fourth Normal Form
- Overall Database Design Process





First Normal Form

- Domain is **atomic** if its elements are considered to be indivisible units
 - ❖ Examples of non-atomic domains:
 - ▀ Set of names, composite attributes
 - ▀ Identification numbers like CS101 that can be broken up into parts
- A relational schema R is in **first normal form** if the domains of all attributes of R are atomic
- Non-atomic values complicate storage and encourage redundant (repeated) storage of data
 - ❖ E.g. Set of accounts stored with each customer, set of children stored with each person, etc.



First Normal Form (Contd.)

- Atomicity is actually a property of how the elements of the domain are used.
 - ❖ E.g. Strings would normally be considered indivisible
 - ❖ Suppose that students are given roll numbers which are strings of the form CS0012 or EE1127
 - ❖ If the first two characters are extracted to find the department, the domain of roll numbers is not atomic.
 - ❖ Doing so is a bad idea: leads to encoding of information in application program rather than in the database.





Pitfalls in Relational Database Design

- Relational database design requires that we find a “good” collection of relation schemas. A bad design may lead to
 - ❖ Repetition of Information.
 - ❖ Inability to represent certain information.
- Design Goals:
 - ❖ Avoid redundant data
 - ❖ Ensure that relationships among attributes are represented
 - ❖ Facilitate the checking of updates for violation of database integrity constraints.



Example

- Consider the relation schema:
 $Lending\text{-schema} = (branch\text{-name}, branch\text{-city}, assets,$
 $customer\text{-name}, loan\text{-number}, amount)$
- Redundancy:
 - ❖ Data for *branch-name*, *branch-city*, *assets* are repeated for each loan that a branch makes
 - ❖ Wastes space
 - ❖ Complicates updating, introducing possibility of inconsistency of *assets* value
- Null values
 - ❖ Cannot store information about a branch if no loans exist
 - ❖ Can use null values, but they are difficult to handle.



Decomposition

- Decompose the relation schema *Lending-schema* into:

Branch-schema = (*branch-name*, *branch-city*, *assets*)

Loan-info-schema = (*customer-name*, *loan-number*,
branch-name, *amount*)


- All attributes of an original schema (R) must appear in the decomposition (R_1, R_2):

$$R = R_1 \cup R_2$$

- Lossless-join decomposition.

For all possible relations r on schema R

$$r = \Pi_{R1}(r) \bowtie \Pi_{R2}(r)$$



Example of Non Lossless-Join Decomposition

- Decomposition of $R = (A, B)$

$$R_1 = (A) \quad R_2 = (B)$$

A	B
α	1
α	2
β	1

$$\Pi_A(r)$$

A
α
β

$$\Pi_A(r)$$

B
1
2

$$\Pi_B(r)$$

A	B
α	1
α	2
β	1
β	2






Goal — Devise a Theory for the Following

- Decide whether a particular relation R is in “good” form.
- In the case that a relation R is not in “good” form, decompose it into a set of relations $\{R_1, R_2, \dots, R_n\}$ such that
 - ❖ each relation is in good form
 - ❖ the decomposition is a lossless-join decomposition
- Our theory is based on:
 - ❖ functional dependencies
 - ❖ multivalued dependencies



Functional Dependencies

- Constraints on the set of legal relations.
- Require that the value for a certain set of attributes determines uniquely the value for another set of attributes.
- A functional dependency is a generalization of the notion of a key.



Functional Dependencies (Cont.)

- Let R be a relation schema $\alpha \subseteq R$ and $\beta \subseteq R$
- The functional dependency

$$\alpha \rightarrow \beta$$

holds on R if and only if for any legal relations $r(R)$, whenever any two tuples t_1 and t_2 of r agree on the attributes α , they also agree on the attributes β . That is,

$$t_1[\alpha] = t_2[\alpha] \Rightarrow t_1[\beta] = t_2[\beta]$$

- Example: Consider $r(A,B)$ with the following instance of r .

1	4
1	5
3	7

- On this instance, $A \rightarrow B$ does **NOT** hold, but $B \rightarrow A$ does hold.



Functional Dependencies (Cont.)

- K is a superkey for relation schema R if and only if $K \rightarrow R$
- K is a candidate key for R if and only if
 - $K \rightarrow R$, and
 - for no $\alpha \subset K$, $\alpha \rightarrow R$
- Functional dependencies allow us to express constraints that cannot be expressed using superkeys. Consider the schema:

Loan-info-schema = (customer-name, loan-number, branch-name, amount).

We expect this set of functional dependencies to hold:

$loan-number \rightarrow amount$
 $loan-number \rightarrow branch-name$

but would not expect the following to hold:

$loan-number \rightarrow customer-name$



Example

Drinkers(name, addr, beersLiked, manf, favoriteBeer)

name	addr	beersLiked	manf	favoriteBeer
Janeway	Voyager	Bud	A.B.	WickedAle
Janeway	Voyager	WickedAle	Pete's	WickedAle
Spock	Enterprise	Bud	A.B.	Bud

- Reasonable FD's to assert:

1. name → addr
2. name → favoriteBeer
3. beersLiked → manf

- Sometimes, several attributes jointly determine another attribute, although neither does by itself. Example:

beer bar → price



Functional Dependencies

- A functional dependency is **trivial** if it is satisfied by all instances of a relation
 - E.g.
 - $\text{customer-name}, \text{loan-number} \rightarrow \text{customer-name}$
 - $\text{customer-name} \rightarrow \text{customer-name}$
 - In general, $\alpha \rightarrow \beta$ is trivial if $\beta \subseteq \alpha$
- “Nontrivial” = right-side attribute not in left side





Closure of a Set of Functional Dependencies

- Given a set F of functional dependencies, there are certain other functional dependencies that are logically implied by F .
 - ⌚ E.g. If $A \rightarrow B$ and $B \rightarrow C$, then we can infer that $A \rightarrow C$
- The set of all functional dependencies logically implied by F is the *closure* of F .
- We denote the closure of F by F^+ .
- We can find all of F^+ by applying Armstrong's Axioms:
 - ⌚ if $\beta \subseteq \alpha$, then $\alpha \rightarrow \beta$ **(reflexivity)**
 - ⌚ if $\alpha \rightarrow \beta$, then $\gamma \alpha \rightarrow \gamma \beta$ **(augmentation)**
 - ⌚ if $\alpha \rightarrow \beta$, and $\beta \rightarrow \gamma$, then $\alpha \rightarrow \gamma$ **(transitivity)**



Example

- $R = \{A, B, C, G, H, I\}$
 $F = \{ A \rightarrow B$
 $\quad A \rightarrow C$
 $\quad CG \rightarrow H$
 $\quad CG \rightarrow I$
 $\quad B \rightarrow H\}$
- some members of F^+
 - ⌚ $A \rightarrow H$
 $\quad \blacksquare$ by transitivity from $A \rightarrow B$ and $B \rightarrow H$
 - ⌚ $AG \rightarrow I$
 $\quad \blacksquare$ by augmenting $A \rightarrow C$ with G , to get $AG \rightarrow CG$
 and then transitivity with $CG \rightarrow I$
 - ⌚ $CG \rightarrow HI$
 $\quad \blacksquare$ from $CG \rightarrow H$ and $CG \rightarrow I$: “union rule” can be inferred from
 - definition of functional dependencies, or
 - Augmentation of $CG \rightarrow I$ to infer $CG \rightarrow CGI$, augmentation of
 $CG \rightarrow H$ to infer $CGI \rightarrow HI$, and then transitivity





Procedure for Computing F^+

- To compute the closure of a set of functional dependencies F :

```

 $F^+ = F$ 
repeat
  for each functional dependency  $f$  in  $F^+$ 
    apply reflexivity and augmentation rules on  $f$ 
    add the resulting functional dependencies to  $F^+$ 
  for each pair of functional dependencies  $f_1$  and  $f_2$  in  $F^+$ 
    if  $f_1$  and  $f_2$  can be combined using transitivity
      then add the resulting functional dependency to  $F^+$ 
until  $F^+$  does not change any further

```

NOTE: We will see an alternative procedure for this task later



Closure of Functional Dependencies (Cont.)

- We can further simplify manual computation of F^+ by using the following additional rules.
 - If $\alpha \rightarrow \beta$ holds and $\alpha \rightarrow \gamma$ holds, then $\alpha \rightarrow \beta\gamma$ holds (**union**)
 - If $\alpha \rightarrow \beta\gamma$ holds, then $\alpha \rightarrow \beta$ holds and $\alpha \rightarrow \gamma$ holds (**decomposition**)
 - If $\alpha \rightarrow \beta$ holds and $\beta \rightarrow \delta$ holds, then $\alpha \rightarrow \delta$ holds (**pseudotransitivity**)

The above rules can be inferred from Armstrong's axioms.



Closure of Attribute Sets

- Given a set of attributes α , define the ***closure*** of α under F (denoted by α^+) as the set of attributes that are functionally determined by α under F :

$$\alpha \rightarrow \beta \text{ is in } F^+ \Leftrightarrow \beta \subseteq \alpha^+$$

- Algorithm to compute α^+ , the closure of α under F

```

result :=  $\alpha$ ;
while (changes to result) do
  for each  $\beta \rightarrow \gamma$  in  $F$  do
    begin
      if  $\beta \subseteq result$  then result := result  $\cup$   $\gamma$ 
    end
  
```



Example of Attribute Set Closure

- $R = (A, B, C, G, H, I)$
- $F = \{A \rightarrow B$
 $A \rightarrow C$
 $CG \rightarrow H$
 $CG \rightarrow I$
 $B \rightarrow H\}$
- $(AG)^+$
 1. $result = AG$
 2. $result = ABCG$ ($A \rightarrow C$ and $A \rightarrow B$)
 3. $result = ABCGH$ ($CG \rightarrow H$ and $CG \subseteq AGBC$)
 4. $result = ABCGHI$ ($CG \rightarrow I$ and $CG \subseteq AGBCH$)
- Is AG a candidate key?
 1. Is AG a super key?
 1. Does $AG \rightarrow R? == Is(AG)^+ \supseteq R$
 2. Is any subset of AG a superkey?
 1. Does $A \rightarrow R? == Is(A)^+ \supseteq R$
 2. Does $G \rightarrow R? == Is(G)^+ \supseteq R$





Uses of Attribute Closure

There are several uses of the attribute closure algorithm:

- Testing for superkey:
 - ❖ To test if α is a superkey, we compute α^+ , and check if α^+ contains all attributes of R .
- Testing functional dependencies
 - ❖ To check if a functional dependency $\alpha \rightarrow \beta$ holds (or, in other words, is in F^+), just check if $\beta \subseteq \alpha^+$.
 - ❖ That is, we compute α^+ by using attribute closure, and then check if it contains β .
 - ❖ Is a simple and cheap test, and very useful
- Computing closure of F
 - ❖ For each $\gamma \subseteq R$, we find the closure γ^+ , and for each $S \subseteq \gamma^+$, we output a functional dependency $\gamma \rightarrow S$.



Goals of Normalization

- Decide whether a particular relation R is in “good” form.
- In the case that a relation R is not in “good” form, decompose it into a set of relations $\{R_1, R_2, \dots, R_n\}$ such that
 - ❖ each relation is in good form
 - ❖ the decomposition is a lossless-join decomposition
- Our theory is based on:
 - ❖ functional dependencies
 - ❖ multivalued dependencies



Decomposition


- Decompose the relation schema *Lending-schema* into:
 $\text{Branch-schema} = (\text{branch-name}, \text{branch-city}, \text{assets})$
 $\text{Loan-info-schema} = (\text{customer-name}, \text{loan-number}, \text{branch-name}, \text{amount})$
- All attributes of an original schema (R) must appear in the decomposition (R_1, R_2):

$$R = R_1 \cup R_2$$

- Lossless-join decomposition.
 For all possible relations r on schema R

$$r = \prod_{R_1}(r) \bowtie \prod_{R_2}(r)$$

- A decomposition of R into R_1 and R_2 is lossless join if and only if at least one of the following dependencies is in F^+ :
 - ❖ $R_1 \cap R_2 \rightarrow R_1$
 - ❖ $R_1 \cap R_2 \rightarrow R_2$




Normalization Using Functional Dependencies

- When we decompose a relation schema R with a set of functional dependencies F into R_1, R_2, \dots, R_n we want
 - ❖ **Lossless-join decomposition:** Otherwise decomposition would result in information loss.
 - ❖ **No redundancy:** The relations R_i preferably should be in either Boyce-Codd Normal Form or Third Normal Form.
 - ❖ **Dependency preservation:** Let F_i be the set of dependencies F^+ that include only attributes in R_i .
 - ◻ Preferably the decomposition should be **dependency preserving**, that is, $(F_1 \cup F_2 \cup \dots \cup F_n)^+ = F^+$
 - ◻ Otherwise, checking updates for violation of functional dependencies may require computing joins, which is expensive.



Example

- $R = (A, B, C)$
 $F = \{A \rightarrow B, B \rightarrow C\}$
 - ❖ Can be decomposed in two different ways
- $R_1 = (A, B), R_2 = (B, C)$
 - ❖ Lossless-join decomposition:
 $R_1 \cap R_2 = \{B\}$ and $B \rightarrow BC$
 - ❖ Dependency preserving
- $R_1 = (A, B), R_2 = (A, C)$
 - ❖ Lossless-join decomposition:
 $R_1 \cap R_2 = \{A\}$ and $A \rightarrow AB$
 - ❖ Not dependency preserving
 (cannot check $B \rightarrow C$ without computing $R_1 \bowtie R_2$)



Second Normal Form

A relation schema R is in 2NF respect to a set F of functional dependencies if for all nonkey set of attributes β holds:

- $\alpha \rightarrow \beta$ where α is a superkey for R

A relation is said to be in Second Normal Form when every nonkey attribute is fully functionally dependent on the primary key. (**No attribute dependent on a portion of primary key**)

- ❖ That is, every nonkey attribute needs the full primary key for unique identification
- ❖ It is important only in cases of keys containing more than one attribute



Example


Drinkers (name, addr, beersLiked, manf, favoriteBeer)

name	addr	beersLiked	manf	favoriteBeer
Janeway	Voyager	Bud	A.B.	WickedAle
Janeway	Voyager	WickedAle	Pete's	WickedAle
Spock	Enterprise	Bud	A.B.	Bud

FD's: $\text{name} \rightarrow \text{addr}$, $\text{name} \rightarrow \text{favoriteBeer}$, $\text{beersLiked} \rightarrow \text{manf}$ violates 2NF.

Lending-schema (branch-name, branch-city, assets,
customer-name, loan-number, amount)

FD's: $\text{branch-name } \text{branch-city} \rightarrow \text{assets}$ violates 2NF.



Boyce-Codd Normal Form


A relation schema R is in BCNF with respect to a set F of functional dependencies if for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$, at least one of the following holds:

- $\alpha \rightarrow \beta$ is trivial (i.e., $\beta \subseteq \alpha$)
- α is a superkey for R

R is in BCNF if for every nontrivial FD for R , say $X \rightarrow A$, then X is a superkey.

Follow from the idea “key \rightarrow everything.”


1. Guarantees no redundancy due to FD's.
2. Guarantees no *update anomalies* = one occurrence of a fact is updated, not all.
3. Guarantees no *deletion anomalies* = valid fact is lost when tuple is deleted.





Example

- $R = (A, B, C)$
 $F = \{A \rightarrow B$
 $B \rightarrow C\}$
 Key = {A}
- R is not in BCNF
- Decomposition $R_1 = (A, B), R_2 = (B, C)$
 - ❖ R_1 and R_2 in BCNF
 - ❖ Lossless-join decomposition
 - ❖ Dependency preserving



Example of Problems


Drinkers(name, addr, beersLiked, manf, favoriteBeer)

name	addr	beersLiked	manf	favoriteBeer
Janeway	Voyager	Bud	A.B.	WickedAle
Janeway	???	WickedAle	Pete's	???
Spock	Enterprise	Bud	???	Bud

FD's:

1. name \rightarrow addr
2. name \rightarrow favoriteBeer
3. beersLiked \rightarrow manf

- ???'s are redundant, since we can figure them out from the FD's.
- Update anomalies: If Janeway gets transferred to the *Intrepid*, will we change addr in each of her tuples?
- Deletion anomalies: If nobody likes Bud, we lose track of Bud's manufacturer.





Each of the given FD's is a BCNF violation:

- Key = {name, beersLiked}
- ▷ Each of the given FD's has a left side that is a proper subset of the key.

Another Example

Beers(name, manf, manfAddr). (Note: 2NF is satisfied)

- FD's = name \rightarrow manf, manf \rightarrow manfAddr.
- Only key is name.
- ▷ manf \rightarrow manfAddr violates BCNF with a left side unrelated to any key.



Testing for BCNF

- To check if a non-trivial dependency $\alpha \rightarrow \beta$ causes a violation of BCNF
 1. compute α^+ (the attribute closure of α), and
 2. verify that it includes all attributes of R , that is, it is a superkey of R .
- **Simplified test:** To check if a relation schema R is in BCNF, it suffices to check only the dependencies in the given set F for violation of BCNF, rather than checking all dependencies in F^+ .
 - ▷ If none of the dependencies in F causes a violation of BCNF, then none of the dependencies in F^+ will cause a violation of BCNF either.
- However, using only F is **incorrect** when testing a relation in a decomposition of R
 - ▷ E.g. Consider $R(A, B, C, D)$, with $F = \{ A \rightarrow B, B \rightarrow C \}$
 - ▷ Decompose R into $R_1(A, B)$ and $R_2(A, C, D)$
 - ▷ Neither of the dependencies in F contain only attributes from (A, C, D) so we might be misled into thinking R_2 satisfies BCNF.
 - ▷ In fact, dependency $A \rightarrow C$ in F^+ shows R_2 is not in BCNF.





BCNF Decomposition Algorithm (1)

```

result := {R};
done := false;
compute F+;
while (not done) do
  if (there is a schema Ri in result that is not in BCNF)
    then begin
      let α → β be a nontrivial functional
      dependency that holds on Ri
      such that α → Ri is not in F+,
      and α ∩ β = ∅;
      result := (result - Ri) ∪ (Ri - β) ∪ (α, β);
    end
  else done := true;

```

Note: each R_i is in BCNF, and decomposition is lossless-join.



Example of BCNF Decomposition

- $R = (\text{branch-name}, \text{branch-city}, \text{assets}, \text{customer-name}, \text{loan-number}, \text{amount})$
 $F = \{\text{branch-name} \rightarrow \text{assets} \text{ branch-city}$
 $\text{loan-number} \rightarrow \text{amount} \text{ branch-name}\}$
Key = {loan-number, customer-name}
- Decomposition
 - ❖ $R_1 = (\text{branch-name}, \text{branch-city}, \text{assets})$
 - ❖ $R_2 = (\text{branch-name}, \text{customer-name}, \text{loan-number}, \text{amount})$
 - ❖ $R_3 = (\text{branch-name}, \text{loan-number}, \text{amount})$
 - ❖ $R_4 = (\text{customer-name}, \text{loan-number})$
- Final decomposition
 R_1, R_3, R_4





BCNF Decomposition Algorithm (2)


Setting: relation R , given FD's F .

Suppose relation R has BCNF violation $X \rightarrow B$.

1. Compute X^+ .


- ☞ Cannot be all attributes – why?

2. Decompose R into X^+ and $(R - X^+) \cup X$.



3. Find the FD's for the decomposed relations.

- ☞ Project the FD's from F = calculate all consequents of F that involve only attributes from X^+ or only from $(R - X^+) \cup X$.



Example

$R = \text{Drinkers}(\underline{\text{name}}, \underline{\text{addr}}, \underline{\text{beersLiked}}, \underline{\text{manf}}, \underline{\text{favoriteBeer}})$

$F =$

1. $\text{name} \rightarrow \text{addr}$
2. $\text{name} \rightarrow \text{favoriteBeer}$
3. $\text{beersLiked} \rightarrow \text{manf}$

Pick BCNF violation $\text{name} \rightarrow \text{addr}$.

- Close the left side: $\text{name}^+ = \text{name } \text{addr } \text{favoriteBeer}$.
- Decomposed relations:
 - $\text{Drinkers1}(\underline{\text{name}}, \underline{\text{addr}}, \underline{\text{favoriteBeer}})$
 - $\text{Drinkers2}(\underline{\text{name}}, \underline{\text{beersLiked}}, \underline{\text{manf}})$
- Projected FD's (skipping a lot of work that leads nowhere interesting):
 - ☞ For Drinkers1: $\text{name} \rightarrow \text{addr}$ and $\text{name} \rightarrow \text{favoriteBeer}$.
 - ☞ For Drinkers2: $\text{beersLiked} \rightarrow \text{manf}$.





(Repeating)

- Decomposed relations:
 - Drinkers1(name, addr, favoriteBeer)
 - Drinkers2(name, beersLiked, manf)
- Projected FD's:
 - For Drinkers1: name → addr and name → favoriteBeer.
 - For Drinkers2: beersLiked → manf.
- BCNF violations?
 - For Drinkers1, name is key and all left sides of FD's are superkeys.
 - For Drinkers2, {name, beersLiked} is the key, and beersLiked → manf violates BCNF.



Decompose Drinkers2

- First set of decomposed relations:
 - Drinkers1(name, addr, favoriteBeer)
 - Drinkers2(name, beersLiked, manf)
- Close $\text{beersLiked}^+ = \text{beersLiked}, \text{manf}$.
- Decompose Drinkers2 into:
 - Drinkers3(beersLiked, manf)
 - Drinkers4(name, beersLiked)
- Resulting relations are all in BCNF:
 - Drinkers1(name, addr, favoriteBeer)
 - Drinkers3(beersLiked, manf)
 - Drinkers4(name, beersLiked)





Testing Decomposition for BCNF

- To check if a relation R_i in a decomposition of R is in BCNF,
 - ❖ Either test R_i for BCNF with respect to the **restriction** of F to R_i (that is, all FDs in F^+ that contain only attributes from R_i)
 - ❖ or use the original set of dependencies F that hold on R , but with the following test:
 - for every set of attributes $\alpha \subseteq R_i$, check that α^+ (the attribute closure of α) either includes no attribute of $R - \alpha$, or includes all attributes of R_i .
 - If the condition is violated by some $\alpha \rightarrow \beta$ in F , the dependency $\alpha \rightarrow (\alpha^+ - \alpha) \cap R_i$ can be shown to hold on R_i , and R_i violates BCNF.
 - We use above dependency to decompose R_i



Database System Concepts

7.39

©Silberschatz, Korth and Sudarshan



BCNF and Dependency Preservation

It is not always possible to get a BCNF decomposition that is dependency preserving

- $R = (J, K, L)$
 $F = \{JK \rightarrow L\}$
 $L \rightarrow K\}$
Two candidate keys = JK and JL
- R is not in BCNF
- Any decomposition of R will fail to preserve

$$JK \rightarrow L$$



Database System Concepts

7.40

©Silberschatz, Korth and Sudarshan



Third Normal Form: Motivation

- There are some situations where
 - ⌚ BCNF is not dependency preserving, and
 - ⌚ efficient checking for FD violation on updates is important
- Solution: define a weaker normal form, called Third Normal Form
 - ⌚ Allows some redundancy (with resultant problems; we will see examples later)
 - ⌚ But FDs can be checked on individual relations without computing a join.
 - ⌚ There is always a lossless-join, dependency-preserving decomposition into 3NF.



Example

One FD structure causes problems:

- If you decompose, you can't check all the FD's only in the decomposed relations.
- If you don't decompose, you violate BCNF.

Abstractly: $R = (A, B, C)$, $F = \{AB \rightarrow C, C \rightarrow B\}$

Example: street city \rightarrow zip, zip \rightarrow city.


Keys: $\{A, B\}$ and $\{A, C\}$, but $C \rightarrow B$ has a left side that is not a superkey.

- Suggests decomposition into $\{B, C\}$ and $\{A, C\}$.
 - ⌚ But you can't check the FD: $AB \rightarrow C$ in only these relations (requires a join)
- Equivalent to example in book:

Banker-schema = (branch-name, customer-name, banker-name)

banker-name \rightarrow branch name

branch name customer-name \rightarrow banker-name



Third Normal Form

- A relation schema R is in third normal form (3NF) if for all:


$$\alpha \rightarrow \beta \text{ in } F^+$$

at least one of the following holds:

- ⌚ $\alpha \rightarrow \beta$ is trivial (i.e., $\beta \in \alpha$)
- ⌚ α is a superkey for R
- ⌚ Each attribute A in $\beta - \alpha$ is contained in a candidate key for R .


(NOTE: each attribute may be in a different candidate key)

- If a relation is in BCNF it is in 3NF (since in BCNF one of the first two conditions above must hold).
- Third condition is a minimal relaxation of BCNF to ensure dependency preservation (will see why later).



Testing for 3NF

- Optimization: Need to check only FDs in F , need not check all FDs in F^+ .
- Use attribute closure to check for each dependency $\alpha \rightarrow \beta$, if α is a superkey.
- If α is not a superkey, we have to verify if each attribute in β is contained in a candidate key of R
 - ⌚ this test is rather more expensive, since it involve finding candidate keys
 - ⌚ testing for 3NF has been shown to be NP-hard
 - ⌚ Interestingly, decomposition into third normal form (described shortly) can be done in polynomial time



3NF Decomposition Algorithm

Let F_c be a canonical cover for F ;
 $i := 0$;
for each functional dependency $\alpha \rightarrow \beta$ in F_c **do**
 if none of the schemas R_j , $1 \leq j \leq i$ contains $\alpha \beta$
 then begin
 $i := i + 1$;
 $R_i := \alpha \beta$
 end
 if none of the schemas R_j , $1 \leq j \leq i$ contains a candidate key for R
 then begin
 $i := i + 1$;
 $R_i :=$ any candidate key for R ;
 end
return (R_1, R_2, \dots, R_i)



What 3NF Gives You

There are two important properties of a decomposition:

1. We should be able to recover from the decomposed relations the data of the original.
 - ❖ Recovery involves projection and join.
2. We should be able to check that the FD's for the original relation are satisfied by checking the projections of those FD's in the decomposed relations.
 - Without proof, we assert that it is always possible to decompose into BCNF and satisfy (1).
 - Also without proof, we can decompose into 3NF and satisfy both (1) and (2).
 - But it is not possible to decompose into BNCF and get both (1) and (2).
 - ❖ Street-city-zip is an example of this point.



Example

- Relation schema:

*Banker-info-schema = (branch-name, customer-name,
banker-name, office-number)*

- The functional dependencies for this relation schema are:

*banker-name → branch-name office-number
customer-name branch-name → banker-name*

- The key is:

{customer-name, branch-name}



Applying 3NF to Banker-info-schema

- The **for** loop in the algorithm causes us to include the following schemas in our decomposition:

*Banker-office-schema = (banker-name, branch-name,
office-number)*

*Banker-schema = (customer-name, branch-name,
banker-name)*


- Since *Banker-schema* contains a candidate key for *Banker-info-schema*, we are done with the decomposition process.





Comparison of BCNF and 3NF

- It is always possible to decompose a relation into relations in 3NF and
 - ❖ the decomposition is lossless
 - ❖ the dependencies are preserved
- It is always possible to decompose a relation into relations in BCNF and
 - ❖ the decomposition is lossless
 - ❖ it may not be possible to preserve dependencies.



Comparison of BCNF and 3NF (Cont.)

- Example of problems due to redundancy in 3NF

$R = (A, B, C)$
 $F = \{AB \rightarrow C, C \rightarrow B\}$

A	B	C
a_1	b_1	c_1
a_2	b_1	c_1
a_3	b_1	c_1
null	b_2	c_2

A schema that is in 3NF but not in BCNF has the problems of


- repetition of information (e.g., the relationship b_1, c_1)
- need to use null values (e.g., to represent the relationship b_2, c_2 where there is no corresponding value for A).





Design Goals

- Goal for a relational database design is:
 - ❖ BCNF.
 - ❖ Lossless join.
 - ❖ Dependency preservation.
- If we cannot achieve this, we accept one of
 - ❖ Lack of dependency preservation
 - ❖ Redundancy due to use of 3NF
- Interestingly, SQL does not provide a direct way of specifying functional dependencies other than superkeys.
Can specify FDs using assertions, but they are expensive to test
- Even if we had a dependency preserving decomposition, using SQL we would not be able to efficiently test a functional dependency whose left hand side is not a key.



Multivalued Dependencies

- There are database schemas in BCNF that do not seem to be sufficiently normalized
- Consider a database

$$\text{classes}(\text{course}, \text{teacher}, \text{book})$$

such that $(c, t, b) \in \text{classes}$ means that t is qualified to teach c , and b is a required textbook for c
- The database is supposed to list for each course the set of teachers any one of which can be the course's instructor, and the set of books, all of which are required for the course (no matter who teaches it).





Multivalued Dependencies (Cont.)

course	teacher	book
database	Avi	DB Concepts
database	Avi	Ullman
database	Hank	DB Concepts
database	Hank	Ullman
database	Sudarshan	DB Concepts
database	Sudarshan	Ullman
operating systems	Avi	OS Concepts
operating systems	Avi	Shaw
operating systems	Jim	OS Concepts
operating systems	Jim	Shaw

classes

- There are no non-trivial functional dependencies and therefore the relation is in BCNF
- Insertion anomalies – i.e., if Sara is a new teacher that can teach database, two tuples need to be inserted

(database, Sara, DB Concepts)
(database, Sara, Ullman)



Multivalued Dependencies (Cont.)

- Therefore, it is better to decompose *classes* into:

course	teacher
database	Avi
database	Hank
database	Sudarshan
operating systems	Avi
operating systems	Jim

teaches

course	book
database	DB Concepts
database	Ullman
operating systems	OS Concepts
operating systems	Shaw

text


We shall see that these two relations are in Fourth Normal Form (4NF)





Multivalued Dependencies Def.

The *multivalued dependency* $X \twoheadrightarrow Y$ holds in a relation R if whenever we have two tuples of R that agree in all the attributes of X , then we can swap their Y components and get two new tuples that are also in R .



Example (Cont.)

- In our example:

$$\begin{aligned} course &\twoheadrightarrow teacher \\ course &\twoheadrightarrow book \end{aligned}$$
- The above formal definition is supposed to formalize the notion that given a particular value of Y ($course$) it has associated with it a set of values of Z ($teacher$) and a set of values of W ($book$), and these two sets are in some sense independent of each other.
- Note:
 - ❖ If $Y \rightarrow Z$ then $Y \twoheadrightarrow Z$
 - ❖ Indeed we have (in above notation) $Z_1 = Z_2$
The claim follows.





Example

Drinkers(name, addr, phones, beersLiked)

with MVD Name $\rightarrow\rightarrow$ phones. If Drinkers has the two tuples:

name	addr	phones	beersLiked
sue	a	p1	b1
sue	a	p2	b2

it must also have the same tuples with phones components swapped:

name	addr	phones	beersLiked
sue	a	p2	b1
sue	a	p1	b2

Note: we must check this condition for *all* pairs of tuples
that agree on name, not just one pair.



MVD Rules

1. Every FD is an MVD.
 - Because if $X \rightarrow Y$, then swapping Y 's between tuples that agree on X doesn't create new tuples.
 - Example, in Drinkers: name $\rightarrow\rightarrow$ addr.
2. *Complementation*: if $X \rightarrow\rightarrow Y$, then $X \rightarrow\rightarrow Z$, where Z is all attributes not in X or Y .
 - Example: since name $\rightarrow\rightarrow$ phones holds in Drinkers, so does name $\rightarrow\rightarrow$ addr beersLiked.




Fourth Normal Form

- A relation schema R is in 4NF with respect to a set D of functional and multivalued dependencies if for all multivalued dependencies in D^+ of the form $\alpha \twoheadrightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$, at least one of the following hold:
 - ▷ $\alpha \twoheadrightarrow \beta$ is trivial (i.e., $\beta \subseteq \alpha$ or $\alpha \cup \beta = R$)
 - ▷ α is a superkey for schema R
- If a relation is in 4NF it is in BCNF

4NF eliminates redundancy due to multiplicative effect of MVD's.

- Formally: R is in Fourth Normal Form if whenever MVD $X \twoheadrightarrow Y$ is *nontrivial* (Y is not a subset of X , and $X \cup Y$ is not all attributes), then X is a superkey.
 - ▷ Remember, $X \rightarrow Y$ implies $X \twoheadrightarrow Y$, so 4NF is more stringent than BCNF.
- Decompose R , using NF violation $X \twoheadrightarrow Y$, into XY and $X \cup (R - Y)$.



Database System Concepts

7.59

©Silberschatz, Korth and Sudarshan

4NF Decomposition Algorithm

```

result := {R};
done := false;
compute D+;
Let Di denote the restriction of D+ to Ri
while (not done)
  if (there is a schema Ri in result that is not in 4NF) then
    begin
      let α →→ β be a nontrivial multivalued dependency that holds
      on Ri such that α → Ri is not in Di, and α ∩ β = φ;
      result := (result - Ri) ∪ (Ri - β) ∪ (α, β);
    end
  else done := true;

```

Note: each R_i is in 4NF, and decomposition is lossless-join



Database System Concepts

7.60

©Silberschatz, Korth and Sudarshan



Splitting Doesn't Hold

Sometimes you need to have several attributes on the right of an MVD. For example:

$\text{Drinkers}(\text{name}, \text{areaCode}, \text{phones}, \text{beersLiked}, \text{beerManf})$

<u>name</u>	<u>areaCode</u>	<u>phones</u>	<u>beersLiked</u>	<u>beerManf</u>
Sue	831	555-1111	Bud	A.B.
Sue	831	555-1111	Wicked Ale	Pete's
Sue	408	555-9999	Bud	A.B.
Sue	408	555-9999	Wicked Ale	Pete's

- $\text{name} \rightarrow\!\!\!}\text{ areaCode}$ phones holds, but neither $\text{name} \rightarrow\!\!\!}\text{ areaCode}$ nor $\text{name} \rightarrow\!\!\!}\text{ phones}$ do.



Example

$\text{Drinkers}(\text{name}, \text{addr}, \text{phones}, \text{beersLiked})$



- FD: $\text{name} \rightarrow \text{addr}$
- Nontrivial MVD's: $\text{name} \rightarrow\!\!\!}\text{ phones}$ and $\text{name} \rightarrow\!\!\!}\text{ beersLiked}$.
- Only key: $\{\text{name}, \text{phones}, \text{beersLiked}\}$
- All three dependencies above violate 4NF.
- Successive decomposition yields 4NF relations:

D1(name, addr)

D2(name, phones)

D3(name, beersLiked)








Overall Database Design Process

- We have assumed schema R is given
 - ❖ R could have been generated when converting E-R diagram to a set of tables.
 - ❖ R could have been a single relation containing *all* attributes that are of interest (called **universal relation**).
 - ❖ Normalization breaks R into smaller relations.
 - ❖ R could have been the result of some ad hoc design of relations, which we then test/convert to normal form.
 - ❖ In practice, usually we start with more relations that intuitively satisfy some normal forms.



ER Model and Normalization

- When an E-R diagram is carefully designed, identifying all entities correctly, the tables generated from the E-R diagram should not need further normalization.
- However, in a real (imperfect) design there can be FDs from non-key attributes of an entity to other attributes of the entity
- E.g. *employee* entity with attributes *department-number* and *department-address*, and an FD $\text{department-number} \rightarrow \text{department-address}$
 - ❖ Good design would have made *department* an entity
- FDs from non-key attributes of a relationship set possible, but rare --- most relationships are binary





Denormalization for Performance


- May want to use non-normalized schema for performance
- E.g. displaying *customer-name* along with *account-number* and *balance* requires join of *account* with *depositor*
- Alternative 1: Use denormalized relation containing attributes of *account* as well as *depositor* with all above attributes
 - ⌚ faster lookup
 - ⌚ Extra space and extra execution time for updates
 - ⌚ extra coding work for programmer and possibility of error in extra code
- Alternative 2: use a materialized view defined as

$$\text{account} \bowtie \text{depositor}$$
 - ⌚ Benefits and drawbacks same as above, except no extra coding work for programmer and avoids possible errors



Other Design Issues

- Some aspects of database design are not caught by normalization
- Examples of bad database design, to be avoided:
 Instead of *earnings(company-id, year, amount)*, use
 - ⌚ *earnings-2000, earnings-2001, earnings-2002*, etc., all on the schema (*company-id, earnings*).
 - ⌚ Above are in BCNF, but make querying across years difficult and needs new table each year
 - ⌚ *company-year(company-id, earnings-2000, earnings-2001, earnings-2002)*
 - ⌚ Also in BCNF, but also makes querying across years difficult and requires new attribute each year.
 - ⌚ Is an example of a **crosstab**, where values for one attribute become column names
 - ⌚ Used in spreadsheets, and in data analysis tools



End of Chapter


Sample *lending* Relation

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>	<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Downtown	Brooklyn	9000000	Jones	L-17	1000
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
Downtown	Brooklyn	9000000	Jackson	L-14	1500
Mianus	Horseneck	400000	Jones	L-93	500
Round Hill	Horseneck	8000000	Turner	L-11	900
Pownal	Bennington	300000	Williams	L-29	1200
North Town	Rye	3700000	Hayes	L-16	1300
Downtown	Brooklyn	9000000	Johnson	L-18	2000
Perryridge	Horseneck	1700000	Glenn	L-25	2500
Brighton	Brooklyn	7100000	Brooks	L-10	2200

Sample Relation r



A	B	C	D
a_1	b_1	c_1	d_1
a_1	b_2	c_1	d_2
a_2	b_2	c_2	d_2
a_2	b_2	c_2	d_3
a_3	b_3	c_2	d_4



Database System Concepts 7.71 ©Silberschatz, Korth and Sudarshan

The customer Relation



<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
Jones	Main	Harrison
Smith	North	Rye
Hayes	Main	Harrison
Curry	North	Rye
Lindsay	Park	Pittsfield
Turner	Putnam	Stamford
Williams	Nassau	Princeton
Adams	Spring	Pittsfield
Johnson	Alma	Palo Alto
Glenn	Sand Hill	Woodside
Brooks	Senator	Brooklyn
Green	Walnut	Stamford



Database System Concepts 7.72 ©Silberschatz, Korth and Sudarshan




The *loan* Relation

<i>loan-number</i>	<i>branch-name</i>	<i>amount</i>
L-17	Downtown	1000
L-23	Redwood	2000
L-15	Perryridge	1500
L-14	Downtown	1500
L-93	Mianus	500
L-11	Round Hill	900
L-29	Pownal	1200
L-16	North Town	1300
L-18	Downtown	2000
L-25	Perryridge	2500
L-10	Brighton	2200



The *branch* Relation

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>
Downtown	Brooklyn	9000000
Redwood	Palo Alto	2100000
Perryridge	Horseneck	1700000
Mianus	Horseneck	400000
Round Hill	Horseneck	8000000
Pownal	Bennington	300000
North Town	Rye	3700000
Brighton	Brooklyn	7100000



The Relation *branch-customer*

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>	<i>customer-name</i>
Downtown	Brooklyn	9000000	Jones
Redwood	Palo Alto	2100000	Smith
Perryridge	Horseneck	1700000	Hayes
Downtown	Brooklyn	9000000	Jackson
Mianus	Horseneck	400000	Jones
Round Hill	Horseneck	8000000	Turner
Pownal	Bennington	300000	Williams
North Town	Rye	3700000	Hayes
Downtown	Brooklyn	9000000	Johnson
Perryridge	Horseneck	1700000	Glenn
Brighton	Brooklyn	7100000	Brooks

Database System Concepts

7.75

©Silberschatz, Korth and Sudarshan

The Relation *customer-loan*

<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Jones	L-17	1000
Smith	L-23	2000
Hayes	L-15	1500
Jackson	L-14	1500
Jones	L-93	500
Turner	L-11	900
Williams	L-29	1200
Hayes	L-16	1300
Johnson	L-18	2000
Glenn	L-25	2500
Brooks	L-10	2200

Database System Concepts


7.76

©Silberschatz, Korth and Sudarshan




The Relation *branch-customer* \bowtie *customer-loan*

<i>branch-name</i>	<i>branch-city</i>	<i>assets</i>	<i>customer-name</i>	<i>loan-number</i>	<i>amount</i>
Downtown	Brooklyn	9000000	Jones	L-17	1000
Downtown	Brooklyn	9000000	Jones	L-93	500
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
Perryridge	Horseneck	1700000	Hayes	L-16	1300
Downtown	Brooklyn	9000000	Jackson	L-14	1500
Mianus	Horseneck	400000	Jones	L-17	1000
Mianus	Horseneck	400000	Jones	L-93	500
Round Hill	Horseneck	8000000	Turner	L-11	900
Pownal	Bennington	300000	Williams	L-29	1200
North Town	Rye	3700000	Hayes	L-15	1500
North Town	Rye	3700000	Hayes	L-16	1300
Downtown	Brooklyn	9000000	Johnson	L-18	2000
Perryridge	Horseneck	1700000	Glenn	L-25	2500
Brighton	Brooklyn	7100000	Brooks	L-10	2200



An Instance of *Banker-schema*


<i>customer-name</i>	<i>banker-name</i>	<i>branch-name</i>
Jones	Johnson	Perryridge
Smith	Johnson	Perryridge
Hayes	Johnson	Perryridge
Jackson	Johnson	Perryridge
Curry	Johnson	Perryridge
Turner	Johnson	Perryridge





Tabular Representation of $\alpha \rightarrow\!\!\!\rightarrow \beta$

	α	β	$R - \alpha - \beta$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$



Relation bc: An Example of Reduncy in a BCNF Relation

<i>loan-number</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
L-23	Smith	North	Rye
L-23	Smith	Main	Manchester
L-93	Curry	Lake	Horseneck




An Illegal bc Relation

<i>loan-number</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
L-23	Smith	North	Rye
L-27	Smith	Main	Manchester

Database System Concepts

7.81

©Silberschatz, Korth and Sudarshan



Decomposition of *loan-info*

<i>branch-name</i>	<i>loan-number</i>
Round Hill	L-58

<i>loan-number</i>	<i>amount</i>

<i>loan-number</i>	<i>customer-name</i>
L-58	Johnson

Database System Concepts

7.82


©Silberschatz, Korth and Sudarshan






Relation of Exercise 7.4

A	B	C
a_1	b_1	c_1
a_1	b_1	c_2
a_2	b_1	c_1
a_2	b_1	c_3





Naracki

Pekot kompletni redosled, vo polmota moze da se koristat dijoker 1, 2, [...] kako kriteriumi za prebaruvanje na bazata (kopceto PRIKAZI)

Cisti **Nova Naracka**
Ovoni Naracka **Brisi Naracka**
Nazad **PRIKAZI**



sdob najdi **Cisti_Click()**
sdel najdi **Prikazi_Click()**
kol najdi **Nov_Click()**
data najdi **Otvori_Click()**

sdob **Brisi_Click()**
sdel **Nazad_Click()**

Form Header
Detail
Form Footer

Öberete go sakaniot zapis, a potoai aktivnosti sto sakate da ja napravite (koristi Ovori Narack za promena, Brisi Narack za brisnenje). Za vnesuvanje novi zapis e kopceto Nova Naracka.

Database System Concepts 1.2 ©Silberschatz, Korth and Sudarshan






Procedure for a Query Building

```
Option Compare Database
Option Explicit
Private Sub AddToWhere(FieldValue As Variant, FieldName As String,
    MyCriteria As String, ArgCount As Integer)
    ' Create criteria for WHERE clause.
    If FieldValue <> "" Then
        ' Add "and" if other criterion exists.
        If ArgCount > 0 Then
            MyCriteria = MyCriteria & " and "
        End If
        ' Append criterion to existing criteria.
        ' Enclose FieldValue and asterisk in quotation marks.
        MyCriteria = (MyCriteria & FieldName & " Like " & Chr(39) & Chr(42) &
            FieldValue & Chr(42) & Chr(39))
        ' Increase argument count.
        ArgCount = ArgCount + 1
    End If
End Sub
```

```
Chr(34) = "
Chr(39) =
Chr(42) = *
```




Procedure for Field Cleaning in a Form

```
Private Sub Cisti_Click()
    Dim MySQL As String
    Dim Tmp As Variant
    MySQL = "SELECT * FROM NajdiNaracka WHERE False"
    ' Clear search text boxes.
    Me![sdob najdi] = Null
    Me![sdel najdi] = Null
    Me![kol najdi] = Null
    Me![data najdi] = Null
    ' Reset subform's RecordSource property to remove records.
    Me![Naracka subform].Form.RecordSource = MySQL

    ' Move insertion point to Look For Company text box.
    Me![sdob najdi].SetFocus
    ' Exit_Cisti_Click:
    ' Exit Sub

    ' Err_Cisti_Click:
    ' MsgBox "Greska-->" & Err.Description, vbInformation, "Greska"
    ' Resume Exit_Cisti_Click
End Sub
```





Procedure for Searching in Database


```

Private Sub Prikazi_Click()
On Error GoTo Err_Prikazi_Click
    Dim MySQL As String, MyCriteria As String, MyRecordSource As String
    Dim ArgCount As Integer
    Dim Tmp As Variant
        ' Initialize argument count.
        ArgCount = 0
        ' Initialize SELECT statement.
        MySQL = "SELECT * FROM NajdiNaracka WHERE "
        MyCriteria = ""
        ' Use values entered in text boxes in form header to create criteria for WHERE clause.
        AddToWhere [sdob najdi], "[NajdiNaracka.sdob]", MyCriteria, ArgCount
        AddToWhere [sdel najdi], "[NajdiNaracka.sdel]", MyCriteria, ArgCount
        AddToWhere [kol najdi], "[NajdiNaracka.kol]", MyCriteria, ArgCount
        AddToWhere [data najdi], "[NajdiNaracka.data]", MyCriteria, ArgCount

        ' If no criterion specified, return all records.
        If MyCriteria = "" Then
            MyCriteria = "True"
        End If

        ' Create SELECT statement.
        MyRecordSource = MySQL & MyCriteria & " ORDER BY sdob"

```



Continues

```


        ' Set RecordSource property of Find Customers Subform.
        Me![Naracka subform].Form.RecordSource = MyRecordSource

        ' If no records match criteria, display message.
        ' Move focus to Clear button.
        If Me![Naracka subform].Form.RecordsetClone.RecordCount = 0 Then
            MsgBox "Nema zapis! ", 48, "Greska"
            Me!Cisti.SetFocus
        Else
            'Enable control in detail section.
            'Me.Section(acDetail).Enabled = True
            'Tmp = EnableControls("Detail", True)
            ' Move insertion point to Find Customers Subform.
            Me![Naracka subform].SetFocus
        End If

        Exit_Prikazi_Click: Exit Sub

        Err_Prikazi_Click:
            MsgBox "Greska-->" & Err.Description, vbInformation, "Greska"
            Resume Exit_Prikazi_Click
    End Sub


```





Do almost Nothing

```
Private Sub Form_Activate()  
    ' Used by Solutions to show toolbar that includes Show Me button.  
    ' Hide built-in Form View toolbar.  
    ' Show Custom Form View toolbar.  
    ' DoCmd.ShowToolbar "Form View", A_TOOLBAR_NO  
    ' DoCmd.ShowToolbar "Custom Form View", A_TOOLBAR_YES  
End Sub  
  
Private Sub Form_Deactivate()  
    ' Used by Solutions to hide toolbar that includes Show Me button.  
    ' Hide Custom Form View toolbar.  
    ' Show built-in Form View toolbar.  
    ' DoCmd.ShowToolbar "Custom Form View", A_TOOLBAR_NO  
    ' DoCmd.ShowToolbar "Form View", A_TOOLBAR_WHERE_APPROP  
End Sub  
  
Private Sub Form_Open(Cancel As Integer)  
    ' Move insertion point to sdob when form is opened.  
    Me![sdob].SetFocus  
End Sub
```



Levels of Abstraction

```
Private Sub Nov_Click()  
On Error GoTo Err_Nov_Click  
Dim stDocName As String  
Dim stLinkCriteria As String  
stDocName = "Naracka"  
DoCmd.OpenForm stDocName, , , stLinkCriteria, acFormAdd  
Exit_Nov_Click: Exit Sub  
Err_Nov_Click:  
    MsgBox Err.Description  
    Resume Exit_Nov_Click  
End Sub  
  
Private Sub Otvori_Click()  
On Error GoTo Err_Otvori_Click  
Dim stDocName As String  
Dim stLinkCriteria As String  
stDocName = "Naracka"  
If IsNull(Me![Naracka subform].Form![sdob]) Then  
    MsgBox "-->Nemate sifra za ovoj zapis!", vbInformation, "Greska"  
Else  
    stLinkCriteria = "[sdob]='" & Me![Naracka subform].Form![sdob] & "'"  
    DoCmd.OpenForm stDocName, , , stLinkCriteria, acFormEdit  
End If  
Exit_Otvori_Click: Exit Sub  
Err_Otvori_Click:  
    MsgBox "Greska-->" & Err.Description, vbInformation, "Greska"  
    Resume Exit_Otvori_Click  
End Sub
```





Instances and Schemas


```
Private Sub Brisi_Click()
On Error GoTo Err_Brisi_Click

Dim MySQL As String, MyCriteria As String

MySQL = "DELETE FROM Naracka WHERE sdob LIKE "
If IsNull(Me![Naracka subform].Form![sdob]) Then
    MsgBox ">->Nemate izbrano zapis!", vbInformation, "Greska"
Else
    MySQL = (MySQL & Chr(34) & Me![Naracka subform].Form![sdob] & Chr(34) & " and sdel
        LIKE " & Chr(34) & Me![Naracka subform].Form![sdel] & Chr(34))
    DoCmd.RunSQL MySQL
    PriKazi_Click
End If

Exit_Brisi_Click: Exit Sub

Err_Brisi_Click:
    MsgBox "Greska-->" & Err.Description, vbInformation, "Greska"
    Resume Exit_Brisi_Click
End Sub
```



Procedure for Opening Form

```
Private Sub Nazad_Click()
On Error GoTo Err_Nazad_Click

Dim stDocName As String
Dim stLinkCriteria As String
DoCmd.Close
stDocName = "GlavnoMeni"
DoCmd.OpenForm stDocName, , , stLinkCriteria

Exit_Nazad_Click: Exit Sub

Err_Nazad_Click:
    MsgBox Err.Description
    Resume Exit_Nazad_Click
End Sub
```

